

Sparkify_Capstone_Project

January 21, 2021

1 Udacity Capstone Project: Sparkify

1.0.1 Project Overview and Problem Statement

This notebook is designed for the Capstone project of Udacity data science nano degree program. We will be working on a subset of digital music streaming service dataset company named Sparkify. It contains details information about existing users and their streaming service usage related behavior. To elaborate how users are using online streaming service, what do they like or not and some more details.

This subset has presented us fact that little over 1/4th =(52) of the total numbers of users 225 churned out of the service throughout the years. This project directs us to find out why and under what circumstances these attritions are happening. The company apparently do not keep any conclusive information from customers to know why they're churning out. So with this dataset we're embarking on a journey mostly analyzing customer behavior, trend and pattern that might lead us finding reasons of churning.

So I will be doing a detail analytical, visual indepth data analysis in search of extracting out why these churning/leaving is happening out of the streaming service. At the onset, intermittently and end I will offer reasons in a data and visualized graphical format must be supported by ingrained data analyses why churning is materializing. The findings will help the company to flesh out plans, promotions and other effective measures to stop churning out and thereby retaining revenue stream flowing.

Here with this notebook I will be working with a smaller subset (128MB) of a larger dataset available (12GB). Assuming this smaller dataset is trully representative of the larger dataset, our analytical steps, ML algorithmic choices can be duplicated or applicable with the larger data set landscape.

First loading the dataset file `mini_sparkify_event_data.json`, thereafter comes all the necessary steps in a sequence to follow through. The programming languages used are Python, PySpark and SQL. Let's start the analytical, logical programming journey...now

1.0.2 Author: Md Ahmed

1.1 Sequential summary of analytical steps:

1. Loading and project related data exploration

- Indepth data exploration with `UserId` relevance:
- 2. Feature definition with categorical and numeric subselection.
- 3. Visualizing the inherent feature-potential with SQL application
 - Six analytical questioning of the dataset
- 4. Page feature distribution analyses with visualization:
 - User and usage analyses of the page features.
- 5. Defining Churn or Attrition.
- 6. Users churning analysis with data attribute facts and visualization.
- 7. Page usage with churning behavior analyses in a questionnaire format
 - Why and how churning is evolving.
- 8. Churning effect on hourly, daily, browser and platforms basis.
- 9. Extensive Feature engineering effect for ML model building.
- 10. Model parameterization performance analyses and summary conclusion.

```
[1]: # import all the needed libraries
from pyspark.sql import SparkSession, Window
from pyspark.sql.functions import udf, col, isnan, when, count, round
from pyspark.sql.functions import avg, stddev, split, udf, isnull, first,
    ↪format_number, rand
from pyspark.sql.functions import sum as Fsum
from pyspark.sql.functions import min as fmin
from pyspark.sql.functions import max as fmax
from pyspark.sql.types import IntegerType, FloatType, LongType
import pyspark.sql.functions as F
from statsmodels.stats.proportion import proportions_ztest

#-----ML
    ↪Packages-----
from pyspark.ml.classification import LogisticRegression, DecisionTreeClassifier
from pyspark.ml.classification import LinearSVC, RandomForestClassifier
from pyspark.ml.feature import StandardScaler, VectorAssembler
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
from pyspark.mllib.evaluation import MulticlassMetrics
from pyspark.ml.tuning import CrossValidator, ParamGridBuilder

from pyspark.ml.feature import Normalizer

from pyspark.ml.tuning import CrossValidatorModel
from pyspark.ml.evaluation import BinaryClassificationEvaluator

%matplotlib inline
```

```
[2]: # import libraries
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pickle as pkl
import seaborn as sns

import os
import re
import copy

from tqdm import *
from time import time
import datetime
```

1.2 Loading and preliminary data exploration

The given small dataset file is `mini_sparkify_event_data.json` was provided by Udacity. We will Load this file and do some basic dataset cleaning and statistical analysis. For instance, finding missing data with or without records, `userIds` or `sessionIds`. Also searching for duplicated data columns and so on...

```
[3]: # create a Spark session
spark = SparkSession.builder.appName('Sparkify_local').getOrCreate()
```

```
[4]: # loading data from source
path = r"C:/Users/paralax11/Desktop/Data_Scientist_Udacity/
↳Capstone_Project_Udacity/mini_sparkify_event_data.json"
df = spark.read.json(path)
```

```
[5]: # Spark data viewing
df.show(3)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+-----+
|          artist|      auth|firstName|gender|itemInSession|lastName|
length|level|          location|method|    page| registration|sessionId|
song|status|          ts|          userAgent|userId|
+-----+-----+-----+-----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+-----+
| Martha Tilston|Logged In|    Colin|    M|          50| Freeman|277.89016|
paid|    Bakersfield, CA|    PUT|NextSong|1538173362000|          29|
Rockpools|    200|1538352117000|Mozilla/5.0 (Wind...|    30|
|Five Iron Frenzy|Logged In|    Micah|    M|          79|    Long|236.09424|
free|Boston-Cambridge-...|    PUT|NextSong|1538331630000|          8|
```

```
Canada| 200|1538352180000|"Mozilla/5.0 (Win...| 9|
| Adam Lambert|Logged In| Colin| M| 51| Freeman| 282.8273|
paid| Bakersfield, CA| PUT|NextSong|1538173362000| 29|Time For
Miracles| 200|1538352394000|Mozilla/5.0 (Wind...| 30|
+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+
only showing top 3 rows
```

```
[6]: df.printSchema()
```

```
root
|-- artist: string (nullable = true)
|-- auth: string (nullable = true)
|-- firstName: string (nullable = true)
|-- gender: string (nullable = true)
|-- itemInSession: long (nullable = true)
|-- lastName: string (nullable = true)
|-- length: double (nullable = true)
|-- level: string (nullable = true)
|-- location: string (nullable = true)
|-- method: string (nullable = true)
|-- page: string (nullable = true)
|-- registration: long (nullable = true)
|-- sessionId: long (nullable = true)
|-- song: string (nullable = true)
|-- status: long (nullable = true)
|-- ts: long (nullable = true)
|-- userAgent: string (nullable = true)
|-- userId: string (nullable = true)
```

1.3 Missing null value analyses with userId:

```
[7]: # Number of null values in different columns
df_null = df.select([count(when(isnan(c) | col(c).isNull(), c)).alias(c) for c_
    ↪in df.columns])
print("Number of Null values inside different columns:")
df_null.show(n=2, truncate=False, vertical=True)
```

```
Number of Null values inside different columns:
-RECORD 0-----
artist      | 58392
auth        | 0
firstName   | 8346
gender      | 8346
```

```

itemInSession | 0
lastName      | 8346
length        | 58392
level         | 0
location      | 8346
method        | 0
page          | 0
registration   | 8346
sessionId     | 0
song          | 58392
status        | 0
ts            | 0
userAgent     | 8346
userId        | 0

```

- We can see that there are 58,392 rows where the artist, length and song attributes are null. That's a large numbers of null to be with the data set. This could be an instance where the artist page was visited by users but they did not play those(artist's) song.
- This is about approx. 20% of the total dataset, a large number indeed and I will explore later...on

We can also see that there are 9 columns with no null values in it: > auth, itemInSession, level, method, page, sessionId, status, ts, userId

A.First on 8,346 null values with 6 columns

```

[182]: # lets find columns with 8,346 number of null values
print("finding columns with total 8,346 numbers of null values:\n")
for cols in df.columns:
    Null_in_columns = df.where(df[cols].isNull()).select(cols).count()
    if Null_in_columns == 8346:
        print(cols)

```

finding columns with total 8,346 numbers of null values:

```

firstName
gender
lastName
location
registration
userAgent

```

We have the same number 8,346 of null values exist **firstName**, **gender**, **lastName**, **location**, **registration** and **userAgent** columns. It is obvious that these null value columns are somehow correlated. Let's explore the possibilities.

- First reason seems that when a user didn't make **registration** then all the other 5 columns stayed empty or null.

- The two other columns `song` and `artist` will also be null because the user didn't play any music.
- Most probably they are guest users.

A clear inter-dependency involves among these columns. To verify the integrated connection I will search little more with these question in mind.

- I think all `registration` must have non-null `firstName`, `gender`, `lastName`, `location` and `userAgent` values in it.
- If an user is not `registered` then the columns must have null-values in those cases.

B.Second on 58,392 null values with 3 columns

```
[183]: # lets find columns with 58392 number of null values
print("finding columns with total of 58,392 numbers of null values:\n")
for cols in df.columns:
    Null_count = df.where(df[cols].isNull()).select(cols).count()
    if Null_count == 58392:
        print(cols)
```

finding columns with total of 58,392 numbers of null values:

```
artist
length
song
```

These `artist`, `length`, `song` columns have altogether 58,392 null values. I can make some logical assumptions as follows:

- When an users was on 'NextSong' page but didn't click on the song to play, then the song column remained null.
- So goes with the `artist` and `length` columns, since user's didn't use those attributes.
- These null values probably stemmed from the guest users who were not registered with the system.
- It seems that `userId` has empty index-value because when users' were allowed to use the system as a 'Guests', which numbered to 8346 of users. So, we will see more evidence that how guest users shaped the dataset.

Registration entanglement:

We can see that if an user is not registered then all these columns like `firstName`, `lastName`, `artist`, `song`.. stays null. On the contray, if `registration` is not-null then these columns contains value.

```
[184]: Name_Reg = df.filter("registration IS NULL AND firstName IS NOT NULL")
print("While registration is null then 'fistName' stays null: ", Name_Reg.
      ↪count())
```

While registration is null then 'fistName' stays null: 0

```
[185]: Location_Reg = df.filter("registration IS NULL AND location IS NOT NULL")
print("While registration is null then 'location' stays null: ", Location_Reg.
      ↪count())
```

While registration is null then 'location' stays null: 0

```
[186]: artist_Reg = df.filter("registration IS NULL AND artist IS NOT NULL")
print("While registration is null then 'artist' column also stays null: ",
      ↪artist_Reg.count())
```

While registration is null then 'artist' column also stays null: 0

```
[187]: song_Reg = df.filter("registration IS NOT NULL AND artist IS NULL")
print("While registration is not-null then 'song' column has value in it: ",
      ↪song_Reg.count())
```

While registration is not-null then 'song' column has value in it: 50046

C. UserId and guests interrelations and "missing/null" values:

Users/ customers are the main subject with this dataset, so I will be doing little indepth understanding about existing users and how they fit into the analyses.

```
[188]: # A simple view of user-situation in the dataset
#df.select(F.countDistinct("userId")).show()
df.describe('userId').show()
print('\nNumber of distinct users in the dataset: ', df.select('userId').
      ↪distinct().count())
```

```
+-----+-----+
|summary|      userId|
+-----+-----+
|  count|      286500|
|   mean| 59682.02278593872|
| stddev|109091.94999910559|
|    min|              |
|    max|              99|
+-----+-----+
```

Number of distinct users in the dataset: 226

- We have 226 distinct users but a total of 286,500 user Ids in the column. So lets check as if there any repeation or empty userIds in the userId column.

```
[11]: #Find out the different categories of the column 'auth' with 'userId' connection
df.select('auth').groupby('auth').count().collect()
```

```
[11]: [Row(auth='Logged Out', count=8249),
      Row(auth='Cancelled', count=52),
      Row(auth='Guest', count=97),
      Row(auth='Logged In', count=278102)]
```

```
[12]: # identifying guest > users who have logged in
guest_logged_in = df.select(['userId', 'auth']).where(col("userId")=='').\
                    where(col('auth').isin(['Guest', 'Logged In'])).count()

print('Number of guest-user with login:', guest_logged_in)
```

Number of guest-user with login: 97

```
[13]: guest_logged_out = df.select("userId", 'auth').where(col("userId")=='').\
        where(col('auth').isin(['Guest', 'Logged Out'])).
        ↪count()

print('Number of guest who logged out:', guest_logged_out)
```

Number of guest who logged out: 8346

GUEST RELATED MISSING VALUES: - Looks like the only empty userIds are where the users are either Logged out or simply guests who have not registered? - So we can see that [auth='Guest', count= 97 and auth='Logged Out', count= 8249] counts total 8346 numbers of users. - It's obvious that these are just not invalid row-data and they have correlation with users being logged out without registration. These are guests users. - So I think we should not delete these unregistered guest-user related rows-columns.

1.3.1 Finding distinct userIds:

```
[14]: df_cleaned = df.filter(df['userId'] != "")
print('Total distinct user_id rows in the dataset: ', df_cleaned.count())
```

Total distinct user_id rows in the dataset: 278154

```
[15]: # check for missing values again after removing empty userId
df_cleaned.select([count(when(isnan(c) | col(c).isNull(), c)) \
                    .alias(c) for c in df.columns]).collect()

print('After final cleaning total uid rows count is:', df_cleaned.count())

# finding distinct users
print('\nOur final distinct user numbers: \n')
df_cleaned.select(F.countDistinct("userId")).show()
```

After final cleaning total uid rows count is: 278154

Our final distinct user numbers:

```
+-----+
|count(DISTINCT userId)|
+-----+
|                225|
+-----+
```

Duplicated columns check

```
[16]: def getDuplicateColumns(df):
        '''
        Get a list of duplicate columns.
        It will iterate over all the columns in dataframe and find the columns
        whose contents are duplicate.
        :param df: Dataframe object
        :return: List of columns whose contents are duplicates.
        '''
        duplicateColumnNames = set()
        # Iterate over all the columns in dataframe
        for x in range(df.shape[1]):
            # Select column at xth index.
            col = df.iloc[:, x]
            # Iterate over all the columns in DataFrame from (x+1)th index till end
            for y in range(x + 1, df.shape[1]):
                # Select column at yth index.
                otherCol = df.iloc[:, y]
                # Check if two columns at x & y index are equal
                if col.equals(otherCol):
                    duplicateColumnNames.add(df.columns.values[y])
        return list(duplicateColumnNames)
#-----
# Source 1: at the bottom of the page
```

```
[17]: duplicated_columns = getDuplicateColumns(df.toPandas())
print('Number of duplicated columns: ', duplicated_columns)
```

Number of duplicated columns: []

```
[18]: # Providing a quick view after all the userId cleaning
pd.DataFrame(data=df_cleaned.tail(2), columns=df_cleaned.columns)
```

```
[18]:
```

	artist	auth	firstName	gender	itemInSession	lastName	\
0	None	Logged In	Emilia	F	44	House	
1	Camera Obscura	Logged In	Emilia	F	45	House	

length	level	location	method	page	\
--------	-------	----------	--------	------	---

0	NaN	paid	New York-Newark-Jersey City, NY-NJ-PA	GET	About
1	170.89261	paid	New York-Newark-Jersey City, NY-NJ-PA	PUT	NextSong

	registration	sessionId	song	status	ts \
0	1538336771000	500	None	200	1543622398000
1	1538336771000	500	The Sun On His Back	200	1543622411000

	userAgent	userId
0	Mozilla/5.0 (compatible; MSIE 9.0; Windows NT ...	300011
1	Mozilla/5.0 (compatible; MSIE 9.0; Windows NT ...	300011

1.4 Feature definiton with categoric and numeric subselection:

Since there are no specific file-column description or data dictionary was provided, I decided to make the following careful logical descriptions of the existing column-attributes. These narratives I hope mirrors the attributes hold with these columns will help us get a level of understanding to make our data analysis more intelligible.

Column Descriptions

USER RELATED INFORMATION

- userId: Unique identifier of an user.
- auth: Indicates an user who is guest, logged in or logged out and did not registered.
- firstName: First name of an user.
- lastName: Users Last name.
- gender: If the user is Male or Female.
- registration: Timestamp for when the user first registered for the service.

USERS SESSION COLUMNS

- itemInSession: The order in the session the event occurred.
- artist: The artist name whose song was played.
- song: The title of the song.
- status: HTTP status code. 2xx=Successful, 3xx=Redirection, 4xx=Client Error.
- length: The length of the song in seconds.
- level: If the user had a free or paid subscription service.
- location: Location of users' by City or State.
- method: HTTP method used, can be GET or PUT.

PAGE - RELATED COLUMNS

- page: Page those were surfed by users'

- sessionId: The unique identifier of a session.
- ts: Timestamp of the user-page-event.
- userAgent: The environment for the user, example OS and web browser used.

1.5 Visualizing inherent feature-potential with SQL:

If we have to work with large dataset that could strain our laptop/desktop's computing ability with memory issues and may force our local machine to freeze. In that case to perform EDA it is an effective choice to load a small subset of data using SQL.

- In order to do that we need to create a temp view within the spark object using the `createOrReplaceTempView` method.
- Before embarking on the detail journey of finding why customer churning is happening, we will evaluate our existing data resources in a questionnaire format to get a better grasp of how our data might fit into our search efforts. Let's see what we got.

```
[19]: # Create Temp View named > df_local_view
df_cleaned.createOrReplaceTempView("df_local_view")
```

1.5.1 1. User distribution with registraion and cancellation?

```
[20]: Registered_to_Cancellation = spark.sql("""
        SELECT userId,
        MAX(CASE WHEN page = 'Cancellation Confirmation' THEN 1
        ↪ELSE 0 END) AS is_cancelled,
        MAX(CASE WHEN page = 'Logout' THEN 1 ELSE 0 END) AS
        ↪Registered
        FROM df_local_view
        GROUP BY userId
        """).toPandas()
```

```
[21]: Registered_to_Cancellation.head(3)
```

```
[21]:   userId  is_cancelled  Registered
0   100010             0            1
1   200002             0            1
2    125             1            0
```

```
[22]: # Existing customer resources based on churned and non-churned customer numbers
users_guest =
    ↪Registered_to_Cancellation[(Registered_to_Cancellation['is_cancelled'] == 1)
    ↪& \
    ↪
    ↪(Registered_to_Cancellation['Registered'] == 0)][['userId']].count()
```

```

print("Number of users did not register and cancelled: ", users_guest)

users_1 =
↳Registered_to_Cancellation[(Registered_to_Cancellation['is_cancelled'] == 1)
↳& \

↳(Registered_to_Cancellation['Registered'] == 1)][ 'userId'].count()

print("Number of users who registered and cancelled: ", users_1)

users_0 =
↳Registered_to_Cancellation[(Registered_to_Cancellation['is_cancelled'] == 0)
↳& \

↳(Registered_to_Cancellation['Registered'] == 1)][ 'userId'].count()

print("Number of users didn't cancel the service and kept registration: ",
↳users_0)

users_2 =
↳Registered_to_Cancellation[(Registered_to_Cancellation['is_cancelled'] == 0)
↳& \

↳(Registered_to_Cancellation['Registered'] == 0)][ 'userId'].count()

print("Number of users didn't cancel but didn't do registration: ", users_2)

```

Number of users did not register and cancelled: 6
 Number of users who registered and cancelled: 46
 Number of users didn't cancel the service and kept registration: 167
 Number of users didn't cancel but didn't do registration: 6

- So we still have 6 users kept the service without doing any registration....are included in the total of $(167 + 6) = 173$ existing users.

1.5.2 2. Gender distribution in our dataset.

```

[23]: gender_count = spark.sql('''
      SELECT gender,COUNT(DISTINCT userId) AS user_counts
      FROM df_local_view
      GROUP BY gender
      ORDER BY user_counts DESC
    ''')

```


```
print('We have more male participants than women:')
gender_count.show()
```

We have more male participants than women:

```
+-----+-----+
|gender|user_counts|
+-----+-----+
|      M|        121|
|      F|        104|
+-----+-----+
```

```
[24]: fig, ax = plt.subplots(figsize=(14, 5), edgecolor='k', facecolor='lightgrey')
ax = sns.barplot(x='gender', y='user_counts', data = gender_count.toPandas())
#plt.bar(x='gender', height='user_counts', width=0.5,
        ↪color='salmon',data=gender_count.toPandas())

plt.title("Number of Male vs. Female participants", fontsize=14)
plt.xlabel('gender', fontsize=13)
plt.ylabel("Count", fontsize=13)
plt.show()
```



1.5.3 3. How does the itemInSession used up by different users?

```
[25]: df_cleaned.select('ItemInSession').describe().collect()
```

```
[25]: [Row(summary='count', ItemInSession='278154'),
      Row(summary='mean', ItemInSession='114.89918174824018'),
      Row(summary='stddev', ItemInSession='129.85172939948959'),
      Row(summary='min', ItemInSession='0'),
```

```
Row(summary='max', ItemInSession='1321')]
```

```
[26]: df_cleaned.groupBy(['userId', 'level']).agg(F.mean('itemInSession'),F.
      ↪max('itemInSession'), F.count('itemInSession')).show(5)
```

```
+-----+-----+-----+-----+-----+
|userId|level|avg(itemInSession)|max(itemInSession)|count(itemInSession)|
+-----+-----+-----+-----+-----+
|100021| free| 64.66771159874608|          185|          319|
|200021| free|44.287553648068666|          104|          233|
|200001| free|24.930379746835442|           67|          158|
|      6| free|16.547619047619047|           46|           84|
|300022| paid| 42.7219730941704|          137|          446|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

```
[27]: df_cleaned.select(['itemInSession']).agg(F.round(F.mean('itemInSession'),2)).
      ↪show()
```

```
+-----+
|round(avg(itemInSession), 2)|
+-----+
|          114.9|
+-----+
```

item-in-Session: In average the mean itemInSessions was 115 minutes time frame among different users.

1.5.4 4. Average song length by usersId's in seconds?

```
[28]: song_session_average = spark.sql('''
      SELECT userId, ROUND(avg(length),2) AS song_session
      FROM df_local_view
      GROUP BY userId
      ORDER BY song_session DESC
      ''')
      song_session_average.show(3)
```

```
+-----+-----+
|userId|song_session|
+-----+-----+
|200001|      267.77|
|   102|      267.27|
|    63|      266.11|
+-----+-----+
```

only showing top 3 rows

```
[29]: df_cleaned.groupby(['userId', 'level']).agg(F.mean('length'), F.max('length'), F.  
      ↪count('length')).show(4)
```

```
+-----+-----+-----+-----+-----+  
|userId|level|      avg(length)|max(length)|count(length)|  
+-----+-----+-----+-----+-----+  
|100021| free|250.57902452173911|  563.35628|         230|  
|200021| free|247.24942137142867|  591.96036|         175|  
|200001| free|267.76714199999986| 1400.2673|         115|  
|      6| free|256.77808180327867|  655.77751|          61|  
+-----+-----+-----+-----+-----+
```

only showing top 4 rows

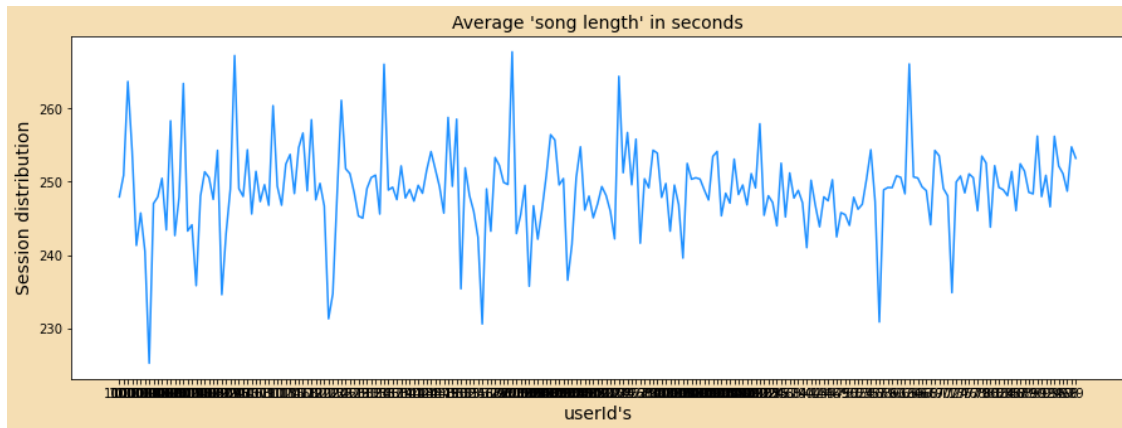
```
[30]: df_cleaned.select(['length']).agg(F.round(F.mean('length'), 2)).show()  
      print('Average length of a song is about ( 249.12/60) = 4.15 minutes ')
```

```
+-----+  
|round(avg(length), 2)|  
+-----+  
|          249.12|  
+-----+
```

Average length of a song is about (249.12/60) = 4.15 minutes

```
[31]: fig, ax = plt.subplots(figsize=(15, 5), edgecolor='k', facecolor='wheat')  
      ax = sns.lineplot(x='userId', y='song_session', color='dodgerblue', data =  
      ↪song_session_average.toPandas())  
  
      plt.title("Average 'song length' in seconds", fontsize=14)  
      plt.xlabel("userId's", fontsize=14)  
      plt.ylabel("Session distribution", fontsize=14)
```

```
[31]: Text(0, 0.5, 'Session distribution')
```



Song lengths: It seems average song length fluctuates in between 230 to 300 seconds.

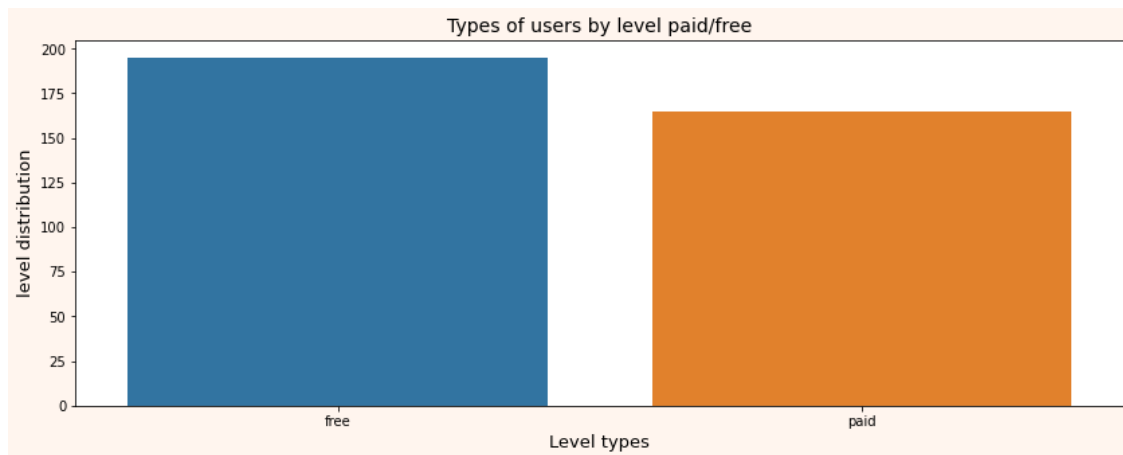
1.5.5 5. Level of subscription type

```
[37]: user_level_counts = spark.sql('''
        SELECT level, COUNT(DISTINCT userId) AS user_level_counts
        FROM df_local_view
        GROUP BY level
        ORDER BY user_level_counts DESC
    ''')
user_level_counts.show()
```

```
+-----+-----+
|level|user_level_counts|
+-----+-----+
| free|              195|
| paid|              165|
+-----+-----+
```

```
[38]: fig, ax = plt.subplots(figsize=(14, 5), edgecolor='k', facecolor='seashell')
ax = sns.barplot(x='level',y='user_level_counts',data=user_level_counts.
    ↳toPandas());
plt.title("Types of users by level paid/free", fontsize=14)
plt.xlabel("Level types", fontsize=13)
plt.ylabel("level distribution", fontsize=13)
```

```
[38]: Text(0, 0.5, 'level distribution')
```

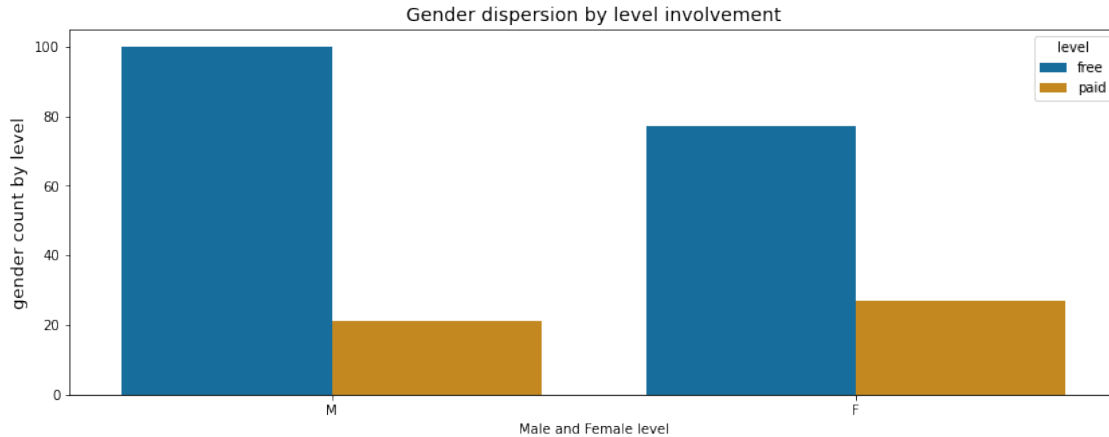



- There are 195 free and 165 paid account in this dataset, in another way, there are 135 users have changed their user-accounts level of subscription.

1.5.6 6. How does the Gender + level distribution looks.

```
[39]: fig, ax = plt.subplots(figsize=(14,5))
gender_level = df.dropDuplicates(['userId']).
    ↳select(['gender','level','userId']).\
                                   groupby('gender','level').count().
    ↳sort('gender', ascending=False).toPandas()

sns.barplot(x='gender', y='count', hue='level', data=gender_level,
    ↳palette='colorblind')
plt.title('Gender dispersion by level involvement', fontsize=14)
plt.xlabel('Male and Female level')
plt.ylabel('gender count by level', fontsize=13)
fig.savefig("gender_level.jpg", bbox_inches='tight')
```



Summary visual: There are more men on free tier than women and conversely more women in paid plan

1.5.7 7. What are the City locations our users mostly coming from?

```
[40]: user_location_count = spark.sql('''
      SELECT location, COUNT(userId) AS user_counts
      FROM df_local_view
      GROUP BY location
      ORDER BY user_counts DESC
      LIMIT 35
      ''').toPandas()
```

```
[41]: user_location_count.head(3)
```

```
[41]:
```

	location	user_counts
0	Los Angeles-Long Beach-Anaheim, CA	30131
1	New York-Newark-Jersey City, NY-NJ-PA	23684
2	Boston-Cambridge-Newton, MA-NH	13873

```
[42]: # split location by city and state
user_location_count = user_location_count.join(user_location_count['location'].\
    ↪str.split(',',expand=True).\
    rename(columns={0:'city',1:'state'})).\
    ↪drop('location',axis=1)
```

```
[43]: user_location_count.head()
```


```
[43]:
```

	user_counts	city	state
0	30131	Los Angeles-Long Beach-Anaheim	CA
1	23684	New York-Newark-Jersey City	NY-NJ-PA

2	13873	Boston-Cambridge-Newton	MA-NH
3	9499	Houston-The Woodlands-Sugar Land	TX
4	7780	Charlotte-Concord-Gastonia	NC-SC

```
[44]: fig, ax = plt.subplots(figsize=(10, 5), edgecolor='k', facecolor='bisque')
ax = user_location_count.sort_values("user_counts", ascending=True).
    ↪set_index("city")["user_counts"].\
        plot(kind="bar",figsize=(17,5), color='lightcoral')

plt.title("Total numbers of users by different cities", fontsize=14)
plt.xlabel("City names", fontsize=12)
plt.xticks(rotation=40)
plt.ylabel("User count distribution", fontsize=12)
plt.grid(True)
```




Obviously, Los Angeles, New York, Boston and Houston are the top cities which have the largest user counts and the rest of the cities have diminished representation.

1.5.8 8. State wise User distribution?

```
[45]: fig, ax = plt.subplots(figsize=(10, 5), edgecolor='k', facecolor='bisque')
ax = user_location_count.sort_values("user_counts", ascending=True).
    ↪set_index("state")["user_counts"].\
        plot(kind="bar",figsize=(17,5), color='steelblue')

plt.title("Total numbers of users by different states", fontsize=14)
plt.xlabel("State names", fontsize=12)
plt.ylabel("User count distribution", fontsize=12)
plt.xticks(rotation=45)
plt.grid(True)
```



We can see that users' are coming from various places from across the united states, but mostly the south-west such as California, Texas.

- The view is partially minimized with 35 city and states. We see that after the 10th city or state the usage trend goes downward consistently not a great trajectory to analyze.

1.5.9 9. Number of Songs played by each user In average and In total

```
[46]: total_song_count = df.select(['userId', 'song']).dropDuplicates().
      ↳groupby(['userId']).count().orderBy('userId')
total_song_count.describe().show()
print('Total song played by all users: ')
```

summary	userId	count
count	226	226
mean	65391.01333333333	894.8141592920354
stddev	105396.4779190716	896.3894099099026
min	1	1
max	99	5947

Total song played by all users:

```
[47]: userId_song = df.select(['userId', 'song']).dropDuplicates().
      ↳groupby(['userId']).count().orderBy('userId')
print('Song played by each user: ', userId_song.show(6))
```

userId	count

			1
	10	630	
	100	2303	
	100001	130	
	100002	194	
	100003	52	

+-----+-----+

only showing top 6 rows

Song played by each user: None

```
[48]: total_song_count = df.select(['userId', 'song']).dropDuplicates().
      ↳groupby(['userId']).count().orderBy('userId')
      print('Average song listened by each user in all session:')
      total_song_count.agg({'count': 'avg'}).
      ↳withColumnRenamed("avg(count)", "mean_song_listened").show()
```

Average song listened by each user in all session:

+-----+-----+
mean_song_listened
+-----+-----+
894.8141592920354
+-----+-----+

Findings: So average user have listened 895 songs in average

1.5.10 10. Number of song played by each user on each session:

```
[49]: song_per_session1 = df.filter(df.page == 'NextSong').
      ↳groupby(['userId', 'sessionId']).count()
      song_per_session2 = song_per_session1.groupby('userId').agg({'count': 'avg'})
      song_per_session2 = song_per_session2.select('userId',
      ↳round(col('avg(count)'), 2))
      song_per_session2 = song_per_session2.withColumnRenamed("round(avg(count), 2)",
      ↳"mean_song_played")
      print("Song played by users in each sessions:")
      song_per_session2.show(5)
```

Song played by users in each sessions:

+-----+-----+
userId mean_song_played
+-----+-----+
100010 39.29
200002 64.5
125 8.0
51 211.1

```
|    124|          145.68|
+-----+-----+
only showing top 5 rows
```

```
[50]: song_per_session2.select('mean_song_played').describe().show()
print('Each users played in average 71 songs in each session:' )
```

```
+-----+-----+
|summary|mean_song_played|
+-----+-----+
|  count|          225|
|   mean|70.7897777777778|
| stddev|42.6154802046791|
|   min|           3.0|
|   max|          286.67|
+-----+-----+
```

Each users played in average 71 songs in each session:

1.5.11 11. Total and average number of artist played by each user:

```
[51]: Num_artist_played_by_session = df.select('userId', 'artist').dropDuplicates().
      ↪groupby('userId').count()
      Num_artist_played_by_session.show(3)
```

```
+-----+-----+
|userId|count|
+-----+-----+
|100010|  253|
|200002|  340|
|   125|    9|
+-----+-----+
```

only showing top 3 rows

```
[52]: print('Descriptive statistics of the artists: ')
      Num_artist_played_by_session.select('count').describe().show()
      print('Each users have played in average 694 artists.')
```

Descriptive statistics of the artists:

```
+-----+-----+
|summary|          count|
+-----+-----+
|  count|          226|
|   mean| 694.29203539823|
| stddev|604.3910901823741|
```

min	1
max	3545

Each users have played in average 694 artists.

1.5.12 12. Song session by each sessionId and by each user on hourly basis:

```
[53]: df_cleaned_copy = df_cleaned.select("*")

[54]: session_end_time = df_cleaned_copy.groupBy('userId', 'sessionId').max('ts').
      ↪withColumnRenamed('max(ts)', 'endTime')
      session_start_time = df_cleaned_copy.groupBy('userId', 'sessionId').min('ts').
      ↪withColumnRenamed('min(ts)', 'starTime')

[55]: # joining both sessions
      session_combined = session_end_time.join(session_start_time, ['userId',
      ↪'sessionId'])
      #session_combined.show(3)

      # 1000 * 60 * 60 = 3600000 milliseconds to 1 hour
      ticks_per_hours = 1000 * 60 * 60 # hourly ticks
      # creating new column
      session_combined = session_combined.withColumn("total_session_time",
      ↪(session_combined.endTime - session_combined.starTime)/(ticks_per_hours))
      # renaming new column
      session_combined = session_combined.select("userId", "sessionId",
      ↪round('total_session_time',2)).\
      ↪
      ↪withColumnRenamed("round(total_session_time, 2)","session_time")
      print("Each user by respective total session and hourly time spend: ")
      session_combined.show(5)
```

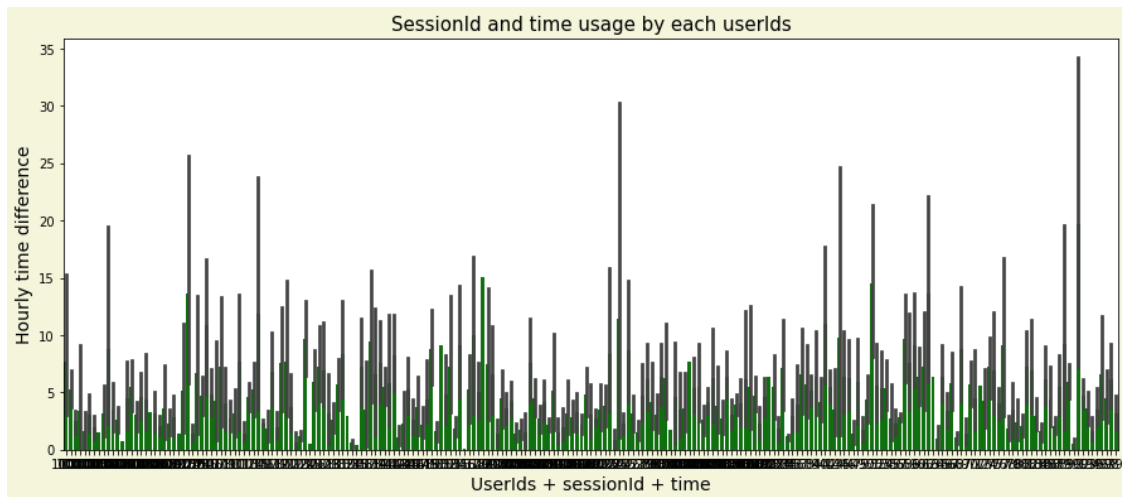
Each user by respective total session and hourly time spend:

userId	sessionId	session_time
101	635	55.0
110	1776	1.39
120	627	15.86
122	691	0.31
140	798	0.16

only showing top 5 rows

```
[56]: fig, ax = plt.subplots(edgecolor='m', facecolor='beige',figsize=(15,6))
sns.barplot(x='userId', y='session_time', data= session_combined.toPandas(),
→color='g')
ax.set_xlabel("UserIds + sessionId + time", fontsize = 14)
ax.set_ylabel("Hourly time difference", fontsize=14)
ax.set_title("SessionId and time usage by each userIds", fontsize=15)
```

```
[56]: Text(0.5, 1.0, 'SessionId and time usage by each userIds')
```



1.5.13 13. Amount of time spent by each user on each session by hours:

```
[57]: df_session = df_cleaned.select("*")
```

```
[58]: # Create a user defined function for formatting the timestamp
get_time = udf(lambda x: datetime.datetime.fromtimestamp(x / 1000.0).
→strftime("%H"))

#Apply the udf on the ts column
df_session = df_session.withColumn("Hours", get_time(df_session.ts))
```

```
[59]: df_session.take(1)
```

```
[59]: [Row(artist='Martha Tilston', auth='Logged In', firstName='Colin', gender='M',
itemInSession=50, lastName='Freeman', length=277.89016, level='paid',
location='Bakersfield, CA', method='PUT', page='NextSong',
registration=1538173362000, sessionId=29, song='Rockpools', status=200,
ts=1538352117000, userAgent='Mozilla/5.0 (Windows NT 6.1; WOW64; rv:31.0)
Gecko/20100101 Firefox/31.0', userId='30', Hours='17')]
```


1.5.14 14.a Total Hour spent by each user

```
[60]: user_total_hour = df_session.select(['userId', 'Hours']).dropDuplicates().\
                                             groupby(['userId']).agg({'Hours':
→ 'sum'})
# renaming 'minutes' to 'session_minutes'
user_total_hour = user_total_hour.withColumnRenamed('sum(Hours)',
→ 'Total_Hours').orderBy('userId')
user_total_hour.show(7)
```

```
+-----+-----+
|userId|Total_Hours|
+-----+-----+
|    10|      276.0|
|   100|      276.0|
|100001|      124.0|
|100002|      169.0|
|100003|       30.0|
|100004|      276.0|
|100005|       93.0|
+-----+-----+
only showing top 7 rows
```

1.5.15 14.b All users with total number of session + total amount of hours with each sessions

```
[61]: # Grouping all users' with total number of sessios + total amount of in all
→ sessions
user_session_hours = df_session.select(['userId', 'sessionId', 'Hours']).
→ dropDuplicates().\
                                             groupby(['userId', 'sessionId']).
→ agg({'Hours': 'sum'})

user_session_hours = user_session_hours.withColumnRenamed('sum(Hours)',
→ 'session_Hours').orderBy('userId')

user_session_hours = user_session_hours.select(['userId', 'sessionId',
→ 'session_Hours']).dropDuplicates().\
                                             groupby('userId').sum().
→ orderBy('userId')
user_session_hours.show(5)
```

```
+-----+-----+-----+
|userId|sum(sessionId)|sum(session_Hours)|
+-----+-----+-----+
```

10	6638	571.0
100	50421	2393.0
100001	151	147.0
100002	496	232.0
100003	89	38.0


+-----+-----+-----+-----+

only showing top 5 rows

```
[62]: print('Shape of this dataset: ', user_session_hours.count(), 'rows and',
        len(user_session_hours.columns), 'columns')
```

Shape of this dataset: 225 rows and 3 columns

```
[63]: fig, ax = plt.subplots(edgecolor='m', facecolor='beige', figsize=(15,6))
sns.barplot(x='userId', y='sum(session_Hours)', data=user_session_hours.
        toPandas(), color='r')
plt.title("Total hours spent by Each user", fontsize=14)
plt.xlabel("Each of 225 userId's", fontsize=13)
plt.ylabel("Hours by each user", fontsize=12)
plt.grid(True)
```



```
[64]: #print((df.count(), len(df.columns)))
#df.count()
#row_number = data.count()
#column_number = len(data.dtypes)
```

1.6 Defining user Churn:

Here we will redefine 'Cancellation Confirmation' page as the churning column, which will contain values '1' for churn and '0' for remained or not churned. We will use this column as a major source of target to investigate out what are the critical reasons of customer churning out of the music streaming service.

We will be subselecting datasets involved with churned data distribution and display critical findings in multiple graphs as we progress through the note book.

```
[65]: # Churn is a label for user who cancelled
cancelled = udf(lambda x: 1 if x == "Cancellation Confirmation" else 0,
    ↳ IntegerType())
# apply to the dataframe
df_cleaned_churned = df_cleaned.withColumn("churn", cancelled(df_cleaned.page))

# Define window bounds
windowval = Window.partitionBy("userId").rangeBetween(Window.
    ↳ unboundedPreceding, Window.unboundedFollowing)
# Applying the window
df_cleaned_churned = df_cleaned_churned.withColumn("churn", Fsum("churn").
    ↳ over(windowval))
```

```
[66]: df_cleaned_churned.take(1)
```

```
[66]: [Row(artist='Sleeping With Sirens', auth='Logged In', firstName='Darianna',
gender='F', itemInSession=0, lastName='Carpenter', length=202.97098,
level='free', location='Bridgeport-Stamford-Norwalk, CT', method='PUT',
page='NextSong', registration=1538016340000, sessionId=31, song='Captain Tyin
Knots VS Mr Walkway (No Way)', status=200, ts=1539003534000,
userAgent='"Mozilla/5.0 (iPhone; CPU iPhone OS 7_1_2 like Mac OS X)
AppleWebKit/537.51.2 (KHTML, like Gecko) Version/7.0 Mobile/11D257
Safari/9537.53"', userId='100010', churn=0)]
```

```
[67]: df_cleaned_churned.dropDuplicates(['userId']).groupby(['churn']).count().show()
```

```
+-----+-----+
|churn|count|
+-----+-----+
|    0|  173|
|    1|   52|
+-----+-----+
```

```
[68]: #df_cleaned_churned.toPandas().drop_duplicates('userId').
    ↳ groupby(['churn'])['churn'].count()
```

1.6.1 Statistical significance test of Churning with Z-test:

source: <https://www.dummies.com/education/math/statistics/how-to-set-up-a-hypothesis-test-null-versus-alternative/>

LEVEL AND CHURNING

- `statsmodels.stats.proportion.proportions_ztest(count, nobs, value=None, alternative='two-sided', prop_var=False)`
- `count` = the number of **successes/target** for each independent sample
- `nobs` (integer or array-like) – the number of **trials/observations**, with the same length as `count`.

Let's set our null hypothesis is "Paid or free service is not effecting users churning behavior" that follows our alternative hypothesis is "Free users were churning more than Paid users".

```
[69]: df_cleaned_churned.dropDuplicates(['userId']).  
      ↪select(['level', 'churn', 'userId']).\  
      ↪groupby('level', 'churn').count().  
      ↪sort('level', ascending=False).collect()
```

```
[69]: [Row(level='paid', churn=0, count=40),  
      Row(level='paid', churn=1, count=8),  
      Row(level='free', churn=0, count=133),  
      Row(level='free', churn=1, count=44)]
```

```
[70]: count = np.array([44, 8])      # the number of successes in number of trials  
      nobs = np.array([177, 48])    # total number of trials or observations
```

```
[71]: stat, pval = proportions_ztest(count, nobs)
```

```
[72]: print('Test static is: {0:0.3f}'.format(stat))
```

Test static is: 1.194

```
[73]: print('P-value is: {0:0.3f}'.format(pval))
```

P-value is: 0.232

Free-Paid service effect on churning:

- Our p-value is $0.23 > 0.05$ is higher than standard 0.05, on which findings we cannot reject the null hypothesis that 'Paid' users are churning more than 'free' users.
- Our Z test-statics is 1.19 which is smaller than 1.96, so we cannot reject the null hypothesis.
- So our null hypothesis is not statistically significant and I can say that **free/paid** services does not influences users churning behavior.

1.7 User Churn effect on features:

- A detail visual analysis of how churning progresses with level an Gender features.

```
[74]: df_panda = df_cleaned_churned.toPandas() # Converting newly changed dataset to pandas dataset for better visual
```

1.7.1 1. Gender and churning spread


```
[75]: # users who churned and used pages
users_1 = df_panda[df_panda.churn == 1].groupby(['gender'])['userId'].count()
users_1 = users_1 / users_1.sum() * 100

# users who didn't cancel but navigated pages
users_0 = df_panda[df_panda.churn == 0].groupby(['gender'])['userId'].count()
users_0 = users_0 / users_0.sum() * 100

# plotting
users_df = pd.DataFrame({'Cancelled service': users_1, 'Remained users': users_0})

#fig, ax = plt.subplots(figsize=(14,5))
ax = users_df.
    ↳ plot(kind='bar', color=('sandybrown', 'forestgreen'), figsize=(14,5));
ax.set_facecolor('ivory')
ax.set_title('Percent of Men and Women cancelled the service(%)', fontsize = 14);
ax.set_ylabel('Percentage of distribution(%)', fontsize = 13);
ax.set_xlabel('Gender distribution on cancellation', fontsize = 14);
ax.set_xticklabels(['Female', 'Male'], rotation=0);

fig.savefig("gender_based_chunred.jpg", bbox_inches='tight')
```



1.7.2 2. Gender + level + churning in detail

```
[76]: # Male-Female churned Paid and Free level
df_cleaned_churned.dropDuplicates(['userId']).groupby(['gender', 'churn']).
    ↪count().sort('gender', ascending=False).show()
```


```
+-----+-----+-----+
|gender|churn|count|
+-----+-----+-----+
|      M|    1|   32|
|      M|    0|   89|
|      F|    1|   20|
|      F|    0|   84|
+-----+-----+-----+
```

```
[77]: # More detail in paid and free service...
df_cleaned_churned.dropDuplicates(['userId']).
    ↪groupby(['gender', 'level', 'churn']).count().sort('level', ascending=False).
    ↪show()
```

```
+-----+-----+-----+-----+
|gender|level|churn|count|
+-----+-----+-----+-----+
|      F| paid|    1|    3|
|      M| paid|    1|    5|
|      F| paid|    0|   24|
|      M| paid|    0|   16|
|      F| free|    1|   17|
|      F| free|    0|   60|
|      M| free|    0|   73|
|      M| free|    1|   27|
+-----+-----+-----+-----+
```

```
[78]: ax = df_cleaned_churned.toPandas().drop_duplicates(['userId']).
    ↪groupby(['gender', 'level', 'churn'])['churn'].count().\
        plot(kind='barh', figsize=(13,6),
    ↪color=('peru', 'seagreen'))

ax.set_facecolor("ivory")
plt.title('Churned is "1" and Remained is "0" by gender and level',fontsize=14)
plt.xlabel('Individual user counts', fontsize=14)
plt.ylabel('gender - level - churn', fontsize=14)
plt.grid()
```



Summary analysis :

We can see that more Male-Free customers churned than Female-free customers. On the contrary, more Male-Paid users churned than Women-Paid customers. Overall in both Paid-Free customer category more Male customer unsubscribe than Female.

1.7.3 3. Percentile projection of churning with gender and level involved

```
[79]: users_1 = df_panda[df_panda.churn == 1].groupby(['gender', 'level'])['userId'].
      ↪count().sort_values(ascending=True)
      # calculating percentage of users
      users_1 = users_1 / users_1.sum()*100

      users_0 = df_panda[df_panda.churn == 0].groupby(['gender', 'level'])['userId'].
      ↪count().sort_values(ascending=True)
      # percentage of users
      users_0 = users_0 / users_0.sum()*100

      users_df = pd.DataFrame({'Cancelled users': users_1, 'Active users': users_0})


      # plotting elements
      ax = users_df.plot(kind='bar', figsize=(14,5),
      ↪color=('sandybrown', 'forestgreen'));
      ax.set_facecolor('ivory')

      ax.set_title('Percentile Female and Male, level based churning or stayed(%)',
      ↪fontsize=14);
      ax.set_xlabel('Male and Female distribution', fontsize=14)
```

```
plt.xticks(rotation=0, horizontalalignment='center', fontweight='light',
    ↳ fontsize=14)
ax.set_xticklabels(['Female-Free', 'Female-Paid', 'Male-Free', 'Male-Paid'],
    ↳ rotation=0);

ax.set_ylabel('Percentile occurrence (%)', fontsize=13)
```

[79]: Text(0, 0.5, 'Percentile occurrence (%)')



1.7.4 4. Non-Percentile Gender and level distribution with churning behavior

```
[80]: # gender and level with churning is true
users_1 = df_panda[df_panda.churn == 1].groupby(['gender', 'level'])['userId'].
    ↳ count().sort_values(ascending=True)
# gender and level with churning not true
users_0 = df_panda[df_panda.churn == 0].groupby(['gender', 'level'])['userId'].
    ↳ count().sort_values(ascending=True)
# inplacing in a data frame
users_df = pd.DataFrame({'Churned': users_1, 'Active users': users_0})


# drawing the plot
ax = users_df.plot(kind='bar', figsize=(14,5));
ax.set_facecolor('seashell')

ax.set_xlabel('Number male-female with free-paid subscriptions', fontsize = 13)
ax.set_ylabel('Number of users', fontsize=12)
ax.set_title('Number of users "paid/free" who cancelled or remained active',
    ↳ fontsize=14)
```



```
plt.xticks(rotation=0, horizontalalignment='right', fontweight='light',  
↪ fontsize=13)
```

[80]: (array([0, 1, 2, 3]), <a list of 4 Text major ticklabel objects>)



Visual analyses: This projection is a non-percentile churning involved with `level + gedner + churning`. To compare I can say that **Percentile visualization** offers a better comparative picture of cancellation and staying with the service than non-percentile one.

1.8 Page usage pattern in details:

Considering this is an online music subscription service, different web page attributes carries users online usage behavior. By analyzing these page-related behavior I will try to extract out a distinct perspective about customer turn overs. Here we will visualize some of the customer subscription related trends, patterns and their effects of cancelling the music service. I used commonsense approach to describe how these pages were designed for this digital service.

In our assigned project we're primarily concerned about what are the reasons existing customers are churning out of the music service. This concern involves a business decision making process. We know by questioning underlying reasons of a business problem we can scrape out better reasoning to solve the problem. In that persuasion I'll be asking these underlying questions about how churning is evolving with this streaming service. These searchings are mainly related to customer behavior patterns and usage trends.

- A. Was that lot of customers were asking for help with number of complaints?
- B. Has any existing customer's usage has dropped drastically lately?
- C. Is 'thumbs down' is a precursor of cancelling the service?
- B. Was there lot of page errors are causing customers to leave the service?
- C. Do existing customer tend to leave the service when their friends unsubscribe?

- D. Paid or free customers who is leaving more in number and why?
- E. Is total music play time is a factor influencing their decision to leave the service?
- F. Adding more song to playlist affects users' decision not to unsubscribe?
- G. Does more thumbsUp means not churning out of the service?

1.8.1 Page descriptions in summary:

Page sub-category attributes descriptions

About: - Information in detail about the song services

Add Friend - Number of friends added while using the service

Add to Playlist - Songs added to users' playlist

Cancel - Users who cancelled the subscription

Cancellation Conf... - Number of users confirmed the cancellation

Downgrade - Number of downgrades by users

Error - Number of error occurred while page surfing

Help - Number of users searched for help

Home - Main page

Logout - How many time users logout of the service page

NextSong - Song lisenting page

Roll Advert - Rolling advertisement while in the page

Save Settings - Saving user settings.

Setting - Page setting

Submit Downgrade - How many users downgraded the service

Submit Upgrade - Users who submit upgrade request

Thumbs Down - Number of Thumbs down or dislikes by users

Thumbs Up - Number of ThumbsUp or likes by users

Upgrade - Page upgrade ny numbers

```
[81]: # Quick view of the web-pages by their page-content and users surfing counts
df_cleaned_churned.select('page').groupby('page').count().
    ↳orderBy('page',ascending=True).show()
```

```
+-----+-----+
|                page| count|
+-----+-----+
|                About|    495|
```

	Add Friend	4277
	Add to Playlist	6526
	Cancel	52
	Cancellation Conf...	52
	Downgrade	2055
	Error	252
	Help	1454
	Home	10082
	Logout	3226
	NextSong	228108
	Roll Advert	3933
	Save Settings	310
	Settings	1514
	Submit Downgrade	63
	Submit Upgrade	159
	Thumbs Down	2546
	Thumbs Up	12551
	Upgrade	499
+-----+		

1.9 Page feature distribution and visualization Snapshot:

Page is a very consequential column in analyzing the whole PySpark project. I will do a detail analyses in a bit later down. But here I'm just introducing all the available page features to its users.

1.9.1 1. Page distribution by page_usage_counts

```
[82]: page_usage_counts = spark.sql('''
        SELECT page,COUNT(userId) AS page_counts
        FROM df_local_view
        GROUP BY page
        ORDER BY page_counts DESC
    ''').toPandas()
page_usage_counts
```

```
[82]:
```


	page	page_counts
0	NextSong	228108
1	Thumbs Up	12551
2	Home	10082
3	Add to Playlist	6526
4	Add Friend	4277
5	Roll Advert	3933
6	Logout	3226
7	Thumbs Down	2546

8	Downgrade	2055
9	Settings	1514
10	Help	1454
11	Upgrade	499
12	About	495
13	Save Settings	310
14	Error	252
15	Submit Upgrade	159
16	Submit Downgrade	63
17	Cancel	52
18	Cancellation Confirmation	52

```
[83]: page_count = df_cleaned.groupby("page").count()

fig, ax = plt.subplots(edgecolor='g', facecolor='bisque')
page_count.toPandas().sort_values(by='count',ascending=True).
    ↪set_index('page')['count'].plot(kind='barh', color='tomato',\

    ↪figsize=(14,6))
plt.title("Page usage counts", fontsize=14)
plt.xlabel("Page distribution count", fontsize=13)
plt.ylabel("Page types", fontsize=12)
plt.grid(True)
```



1.9.2 2. Page distribution by user_counts

```
[84]: page_user_counts = spark.sql('''
      SELECT page,COUNT(DISTINCT userId) AS user_counts
      FROM df_local_view
      GROUP BY page
```

```
ORDER BY user_counts DESC
''').toPandas()


page_user_counts
```

```
[84]:
```

	page	user_counts
0	NextSong	225
1	Home	223
2	Thumbs Up	220
3	Add to Playlist	215
4	Logout	213
5	Roll Advert	207
6	Add Friend	206
7	Thumbs Down	203
8	Settings	195
9	Help	192
10	Upgrade	168
11	About	155
12	Downgrade	154
13	Save Settings	132
14	Submit Upgrade	131
15	Error	117
16	Cancel	52
17	Cancellation Confirmation	52
18	Submit Downgrade	49

```
[85]: fig, ax = plt.subplots(edgecolor='g', facecolor='bisque')
page_user_counts.sort_values('user_counts').set_index("page")["user_counts"].\
plot(kind='barh',
      color='lightseagreen',figsize=(13, 5))
plt.title("Pages usage by number of users", fontsize=14)
plt.xlabel("User counts", fontsize=12)
plt.ylabel("Page index", fontsize=12)
```

```
[85]: Text(0, 0.5, 'Page index')
```



```
[86]: print('Visualizing page by user_counts and usage_counts: \n')
pd.merge(page_user_counts, page_uage_counts, on=['page'], how='inner',
         ↪validate='one_to_one')
```

Visualizing page by user_counts and usage_counts:

```
[86]:
```

	page	user_counts	page_counts
0	NextSong	225	228108
1	Home	223	10082
2	Thumbs Up	220	12551
3	Add to Playlist	215	6526
4	Logout	213	3226
5	Roll Advert	207	3933
6	Add Friend	206	4277
7	Thumbs Down	203	2546
8	Settings	195	1514
9	Help	192	1454
10	Upgrade	168	499
11	About	155	495
12	Downgrade	154	2055
13	Save Settings	132	310
14	Submit Upgrade	131	159
15	Error	117	252
16	Cancel	52	52
17	Cancellation Confirmation	52	52
18	Submit Downgrade	49	63

PAGE USAGE AND SNAPSHOT ANALYSES:

- We can see that out of roughly 225 registered users of which 52 of them end up cancelling the service and around 50 downgraded their service.
- We can see from this graph that users tend to use the service mostly for playing music also they visits the home button a lot of times to rate music that they like.
- Home page counts the second largest visits.
- There are lots of thumbsUp, Add Friend, Add to Playlist, Roll Advert usage is visual.
- There is a large number of 'downgrade' is also present.
- We can see that cancel and cancellation confirmations are the same numbers and meanings.

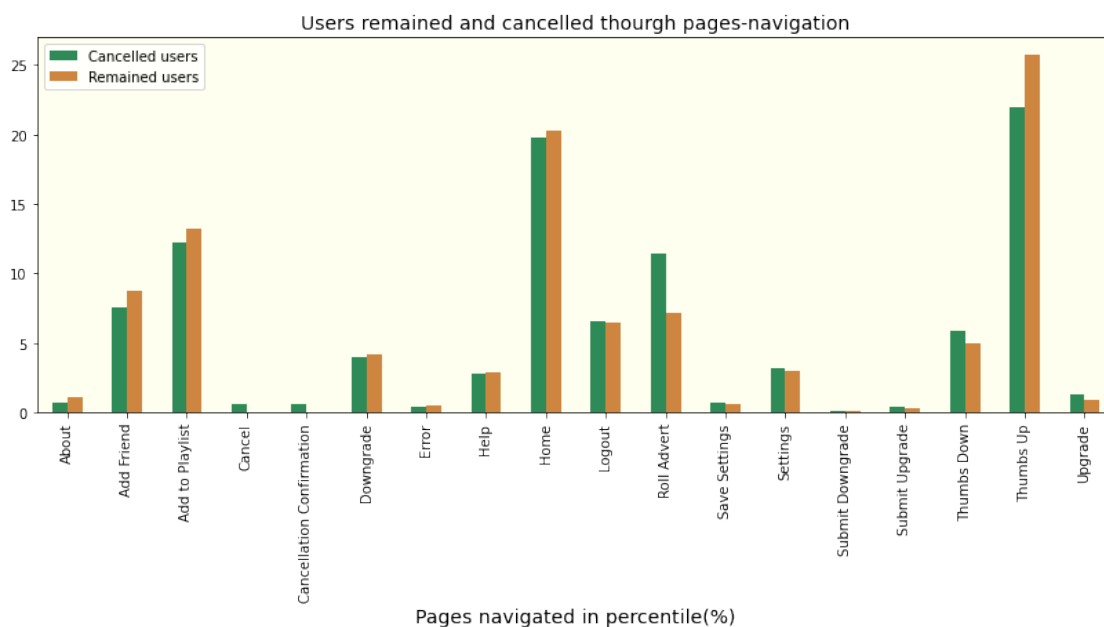
1.10 Page-distribution effect on users churning:

I've dropped NextSong page from total page count, since it has the highest visiting average ($228,108 / 225 = 1013.81$) by each users. It keeps tab everytime an user plays a song, which has the potential of skewing the overall page visiting distribution.

```
[87]: # users who churned and used pages
users_1 = df_panda[df_panda.churn == 1].groupby(['page'])['userId'].count().
↳drop('NextSong').sort_values(ascending=False)
users_1 = users_1 / users_1.sum()*100

# users who didn't cancel but navigated pages
users_0 = df_panda[df_panda.churn == 0].groupby(['page'])['userId'].count().
↳drop('NextSong').sort_values(ascending=False)
users_0 = users_0 / users_0.sum()*100

# plotting
users_df = pd.DataFrame({'Cancelled users': users_1, 'Remained users': users_0})
ax = users_df.plot(kind='bar', figsize=(14,5), color=("seagreen", "peru"));
ax.set_facecolor('ivory')
ax.set_xlabel('Pages navigated in percentile(%)', fontsize = 14);
ax.set_title('Users remained and cancelled through pages-navigation', fontsize_
↳= 14);
```



We can see that pages....'Roll Advert', 'Thumbs Down', 'Upgrade' and 'Settings' are causing more churning out than other pages.

It is obvious that pages like Add Friend, Add to Playlist, Downgrade, Home,

Help,ThumbsUp are keeping users with the service more.

1.10.1 Page attribute related churning behavior analyses with 'Scatterplot'

How users left their markings on churning while navigating different pages can be very clearly visual in these 6 different scatterplots. These scatter plots offers true visualization how churning evolved with page usage.

1.10.2 1. Number of friends added in correlation with cancellation.

```
[88]: # Find song long time correlation w/ cancellation
Friends_vs_Cancellation = spark.sql("""
    SELECT userId,
           MAX(CASE WHEN page = 'Cancellation Confirmation' THEN 1 ELSE 0 END) AS
    ↪ is_cancelled,
           SUM(CASE WHEN page = 'Add Friend' THEN 1 ELSE 0 END) AS
    ↪ Num_Friends_Added
    FROM df_local_view
    GROUP BY userId
""").toPandas()
```

```
[89]: Friends_vs_Cancellation.head(3)
```

```
[89]:
```

	userId	is_cancelled	Num_Friends_Added
0	100010	0	4
1	200002	0	4
2	125	1	0

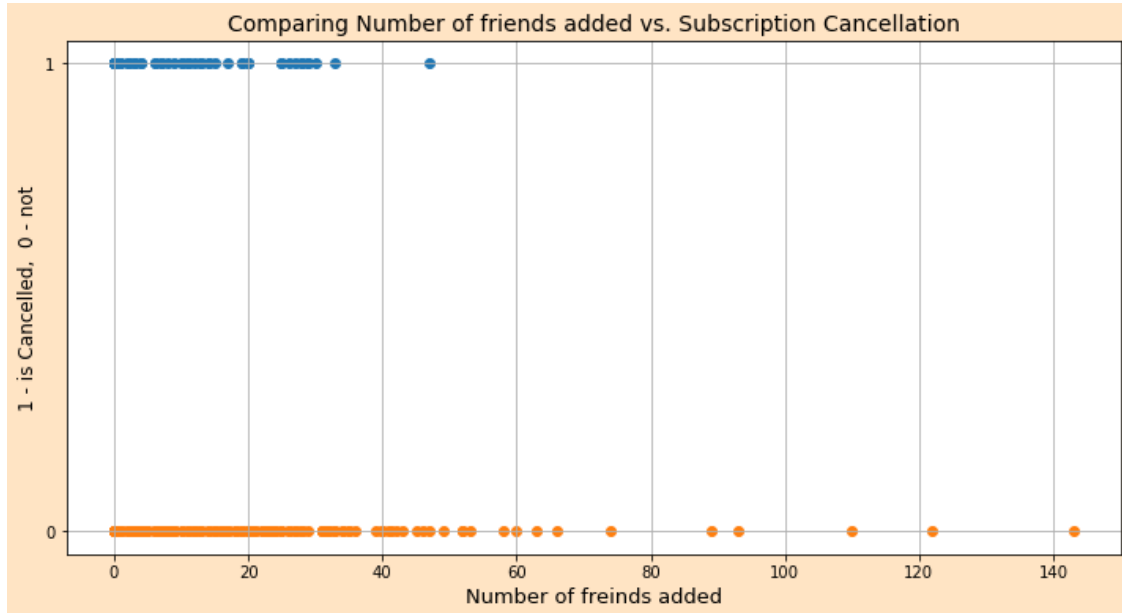
```
[90]: fig, ax = plt.subplots(figsize=(12, 6), edgecolor='k', facecolor='bisque')

plt.scatter(Friends_vs_Cancellation[Friends_vs_Cancellation["is_cancelled"] ==
    ↪ 1]["Num_Friends_Added"],
           Friends_vs_Cancellation[Friends_vs_Cancellation["is_cancelled"] ==
    ↪ 1]["is_cancelled"])
#-----
plt.scatter(Friends_vs_Cancellation[Friends_vs_Cancellation["is_cancelled"] ==
    ↪ 0]["Num_Friends_Added"],
           Friends_vs_Cancellation[Friends_vs_Cancellation["is_cancelled"] ==
    ↪ 0]["is_cancelled"])

plt.yticks((0, 1))
plt.title("Comparing Number of friends added vs. Subscription Cancellation",
    ↪ fontsize = 14)
plt.xlabel("Number of freinds added", fontsize = 13)
plt.ylabel("1 - is Cancelled, 0 - not", fontsize = 12)
```



```
plt.grid()
```



View analysis: We see that the more an user adds friend the the less the chance that they would unsubscribe the service.

1.10.3 2. Add to playlist vs. cancellation.

```
[91]: addPlayList_to_cancellation = spark.sql("""
        SELECT userId,
        MAX(CASE WHEN page = 'Cancellation Confirmation' THEN 1_
        ↪ELSE 0 END) AS is_cancelled,
        SUM(CASE WHEN page = 'Add to Playlist' THEN 1 ELSE 0 END)_
        ↪AS addPlayList
        FROM df_local_view
        GROUP BY userId
    """).toPandas()
```

```
[92]: addPlayList_to_cancellation.head(3)
```

```
[92]:   userId  is_cancelled  addPlayList
0  100010             0             7
1  200002             0             8
2    125             1             0
```

```
[93]: fig, ax = plt.subplots(figsize=(12, 6), edgecolor='b', facecolor='wheat')
```

```

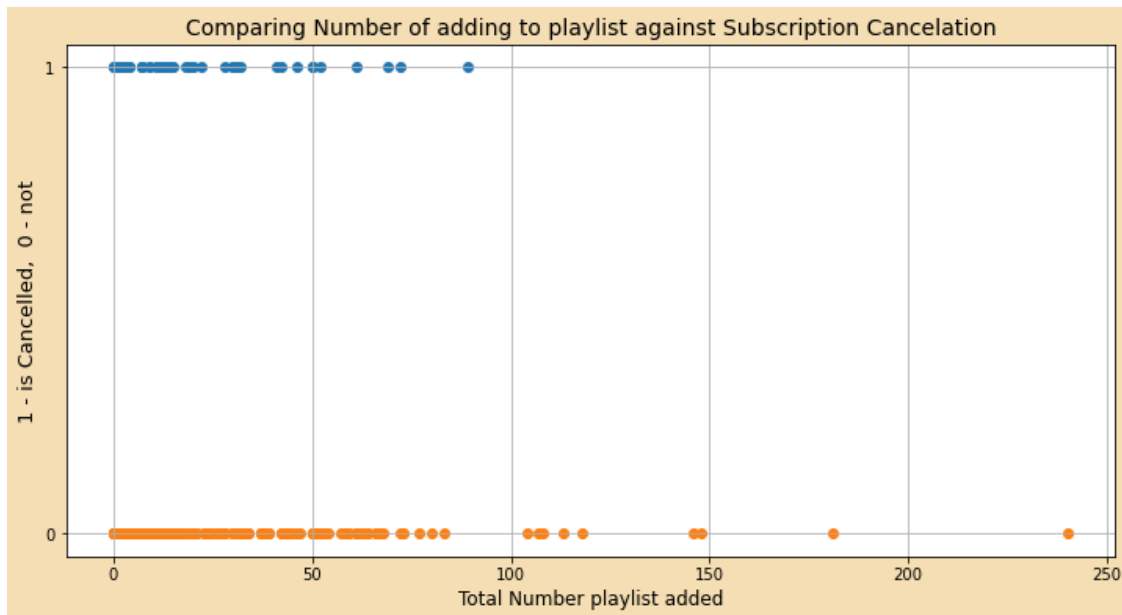
plt.
↳scatter(addPlayList_to_cancellation[addPlayList_to_cancellation["is_cancelled"]
↳== 1] ["addPlayList"],

↳
↳addPlayList_to_cancellation[addPlayList_to_cancellation["is_cancelled"] ==
↳1] ["is_cancelled"])
#-----
plt.
↳scatter(addPlayList_to_cancellation[addPlayList_to_cancellation["is_cancelled"]
↳== 0] ["addPlayList"],

↳
↳addPlayList_to_cancellation[addPlayList_to_cancellation["is_cancelled"] ==
↳0] ["is_cancelled"])

plt.xticks((0, 1))
plt.title("Comparing Number of adding to playlist against Subscription
↳Cancellation", fontsize=14)
↳
plt.xlabel("Total Number playlist added", fontsize=12)
plt.ylabel("1 - is Cancelled, 0 - not", fontsize=13)
plt.grid()

```



View summary: When an user adding more songs in their playlist they tends not to churning out of the service.

1.10.4 3. A juxtaposed view of Add Friend and Add to Playlist

```
[94]: # Find song time in correlation with cancellation
Friend_vs_Playlist = spark.sql("""
    SELECT userId,
           MAX(CASE WHEN page = 'Cancellation Confirmation' THEN 1 ELSE 0 END) AS
    ↪ is_cancelled,
           SUM(CASE WHEN page = 'Add Friend' THEN 1 ELSE 0 END) AS
    ↪ Total_Friend_Added,
           MIN(CASE WHEN page = 'Cancellation Confirmation' THEN 2 ELSE 3 END) AS
    ↪ also_cancelled,
           SUM(CASE WHEN page = 'Add to Playlist' THEN 1 ELSE 0 END) AS
    ↪ Total_Playlist_Added
    FROM df_local_view
    GROUP BY userId
""").toPandas()
```

```
[95]: Friend_vs_Playlist.head(3)
```

```
[95]:   userId  is_cancelled  Total_Friend_Added  also_cancelled  \
0   100010             0                 4             3
1   200002             0                 4             3
2     125             1                 0             2

   Total_Playlist_Added
0                 7
1                 8
2                 0
```

```
[96]: fig, ax = plt.subplots(figsize=(15, 6), edgecolor='b', facecolor='bisque')


plt.scatter(Friend_vs_Playlist[Friend_vs_Playlist["is_cancelled"] ==
    ↪ 1]["Total_Friend_Added"],
           Friend_vs_Playlist[Friend_vs_Playlist["is_cancelled"] ==
    ↪ 1]["is_cancelled"])

plt.scatter(Friend_vs_Playlist[Friend_vs_Playlist["is_cancelled"] ==
    ↪ 0]["Total_Friend_Added"],
           Friend_vs_Playlist[Friend_vs_Playlist["is_cancelled"] ==
    ↪ 0]["is_cancelled"])
#-----
plt.scatter (Friend_vs_Playlist[Friend_vs_Playlist["also_cancelled"] ==
    ↪ 2]["Total_Playlist_Added"],
           Friend_vs_Playlist[Friend_vs_Playlist["also_cancelled"] ==
    ↪ 2]["also_cancelled"])
```

```
plt.scatter(Friend_vs_Playlist[Friend_vs_Playlist["also_cancelled"] == 0],
            Friend_vs_Playlist[Friend_vs_Playlist["also_cancelled"] == 0],
            Friend_vs_Playlist[Friend_vs_Playlist["also_cancelled"] == 1],
            Friend_vs_Playlist[Friend_vs_Playlist["also_cancelled"] == 2],
            Friend_vs_Playlist[Friend_vs_Playlist["also_cancelled"] == 3])

#-----

plt.yticks((0, 1, 2, 3))
plt.title("1-churned 0-remained Add Friend | 2-churned 3-Remained with Add to Playlist",
          fontsize=14)
plt.xlabel("Add Friend (0,1) and Add to Playlist|(2,3) indicators of subscription",
          fontsize=14)
plt.ylabel("Add Friend and Add to Playlist indicators",
          fontsize=14)
plt.grid()
```



Visual analyses: Looking into line 0 and 3 are lot longer than line 1 and 2 indicates pages Add Friend and Add to Playlist are holding more users with the streaming service.

1.10.5 4. Subscription cancellation with thumbs-up and thumbs-down pages

```
[97]: # Find song time in correlation with cancellation
thumbsDown_vs_thumbsUp = spark.sql("""
    SELECT userId,
           MAX(CASE WHEN page = 'Cancellation Confirmation' THEN 1 ELSE 0 END) AS
    is_cancelled,
           SUM(CASE WHEN page = 'Thumbs Down' THEN 1 ELSE 0 END) AS
    Total_Thumbs_Down,
           MIN(CASE WHEN page = 'Cancellation Confirmation' THEN 2 ELSE 3 END) AS
    also_cancelled,
```

```

        SUM(CASE WHEN page = 'Thumbs Up' THEN 1 ELSE 0 END) AS_
↪Total_Thumbs_Up
        FROM df_local_view
        GROUP BY userId
        """).toPandas()

```

```
[98]: thumbsDown_vs_thumbsUp.head(3)
```

```
[98]:
```

	userId	is_cancelled	Total_Thumbs_Down	also_cancelled	Total_Thumbs_Up
0	100010	0	5	3	17
1	200002	0	6	3	21
2	125	1	0	2	0

```
[99]: fig, ax = plt.subplots(figsize=(15, 6), edgecolor='b', facecolor='bisque')

plt.scatter(thumbsDown_vs_thumbsUp[thumbsDown_vs_thumbsUp["is_cancelled"] ==_
↪1] ["Total_Thumbs_Down"],
            thumbsDown_vs_thumbsUp[thumbsDown_vs_thumbsUp["is_cancelled"] ==_
↪1] ["is_cancelled"])

plt.scatter(thumbsDown_vs_thumbsUp[thumbsDown_vs_thumbsUp["is_cancelled"] ==_
↪0] ["Total_Thumbs_Down"],
            thumbsDown_vs_thumbsUp[thumbsDown_vs_thumbsUp["is_cancelled"] ==_
↪0] ["is_cancelled"])

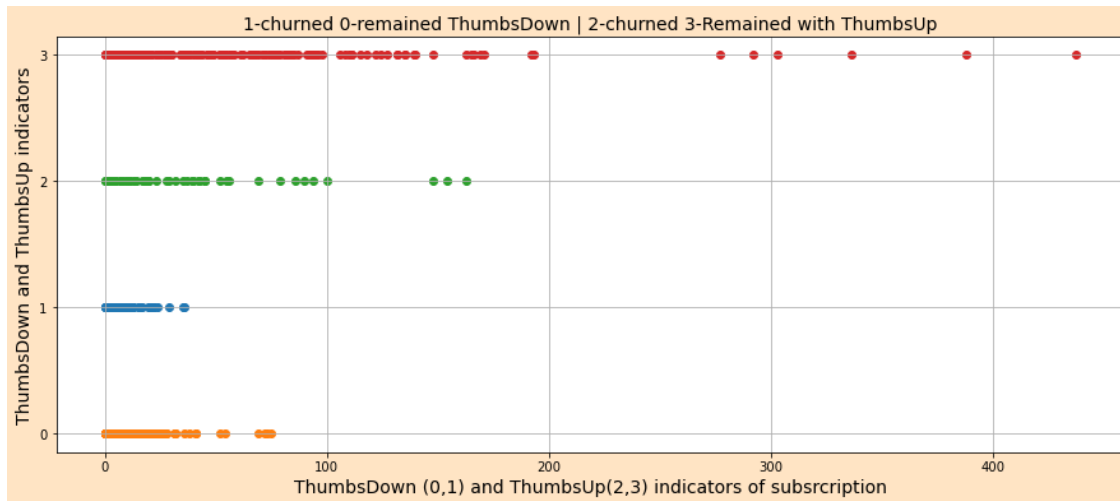
#-----
plt.scatter(thumbsDown_vs_thumbsUp[thumbsDown_vs_thumbsUp["also_cancelled"] ==_
↪2] ["Total_Thumbs_Up"],
            thumbsDown_vs_thumbsUp[thumbsDown_vs_thumbsUp["also_cancelled"] ==_
↪2] ["also_cancelled"])

plt.scatter(thumbsDown_vs_thumbsUp[thumbsDown_vs_thumbsUp["also_cancelled"] ==_
↪3] ["Total_Thumbs_Up"],
            thumbsDown_vs_thumbsUp[thumbsDown_vs_thumbsUp["also_cancelled"] ==_
↪3] ["also_cancelled"])

#-----

plt.yticks((0, 1, 2, 3))
plt.title("1-churned 0-remained ThumbsDown | 2-churned 3-Remained with_
↪ThumbsUp", fontsize=14)
plt.xlabel("ThumbsDown (0,1) and ThumbsUp(2,3) indicators of subsrcription",_
↪fontsize=14)
plt.ylabel("ThumbsDown and ThumbsUp indicators", fontsize=14)
plt.grid()

```



Findings: Obviously thumbsUp lines (2,3) is holding more users with the service however thumbsDownline(0,1) is not causing lot more churning out of the service either. It is not an alarming attrition with the thumbsDown page.

1.10.6 5. Subscription cancellation with page - error and Help

```
[100]: Page_error_help = spark.sql("""
        SELECT userId,
        MAX(CASE WHEN page = 'Cancellation Confirmation' THEN 1 ELSE 0 END) AS is_cancelled,
        SUM(CASE WHEN page = 'Error' THEN 1 ELSE 0 END) AS total_page_error,
        MIN(CASE WHEN page = 'Cancellation Confirmation' THEN 2 ELSE 3 END) AS also_cancelled,
        SUM(CASE WHEN page = 'Help' THEN 1 ELSE 0 END) AS total_help_needed
        FROM df_local_view
        GROUP BY userId
        """).toPandas()
```

```
[101]: Page_error_help.head(3)
```

```
[101]:
```

	userId	is_cancelled	total_page_error	also_cancelled	total_help_needed
0	100010	0	0	3	2
1	200002	0	0	3	2
2	125	1	0	2	0

```
[102]: fig, ax = plt.subplots(figsize=(15, 6), edgecolor='b', facecolor='bisque')
```

```

plt.scatter(Page_error_help[Page_error_help["is_cancelled"] == 1][
    "total_page_error"],
            Page_error_help[Page_error_help["is_cancelled"] == 1][
    "is_cancelled"])

plt.scatter(Page_error_help[Page_error_help["is_cancelled"] == 0][
    "total_page_error"],
            Page_error_help[Page_error_help["is_cancelled"] == 0][
    "is_cancelled"])

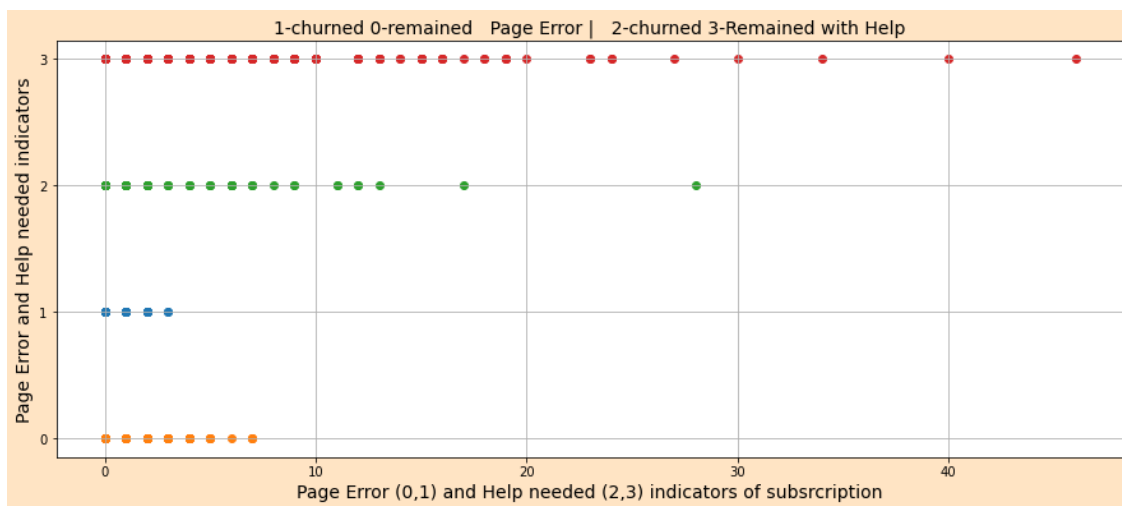
#-----
plt.scatter (Page_error_help[Page_error_help["also_cancelled"] == 2][
    "total_help_needed"],
            Page_error_help[Page_error_help["also_cancelled"] == 2][
    "also_cancelled"])

plt.scatter(Page_error_help[Page_error_help["also_cancelled"] == 3][
    "total_help_needed"],
            Page_error_help[Page_error_help["also_cancelled"] == 3][
    "also_cancelled"])

#-----

plt.yticks((0, 1, 2, 3))
plt.title("1-churned 0-remained Page Error | 2-churned 3-Remained with Help",
    fontsize=14)
plt.xlabel("Page Error (0,1) and Help needed (2,3) indicators of subscription",
    fontsize=14)
plt.ylabel("Page Error and Help needed indicators", fontsize=14)
plt.grid()

```



View summary: We can see there that the number of errors in pages and needed help were relatively lower and users seems not bothered by that. This could just be a coincidence because the sample size is smaller. Help page is keeping more users with the service than page-errors.

1.10.7 6. Rolling Advertisement on unsubscribing

```
[103]: # Find song time in correlation with cancellation
Roll_Advert_cancellation = spark.sql("""
    SELECT userId,
           MAX(CASE WHEN page = 'Cancellation Confirmation' THEN 1 ELSE 0 END) AS
    ↪is_cancelled,
           SUM(CASE WHEN page = 'Roll Advert' THEN 1 ELSE 0 END) AS
    ↪total_roll_advert
    FROM df_local_view
    GROUP BY userId
""").toPandas()
```


```
[104]: Roll_Advert_cancellation.head(4)
```

```
[104]:   userId  is_cancelled  total_roll_advert
0   100010             0                 52
1   200002             0                  7
2     125             1                  1
3      51             1                  0
```

```
[105]: fig, ax = plt.subplots(figsize=(12, 5), edgecolor='b', facecolor='wheat')

plt.scatter(Roll_Advert_cancellation[Roll_Advert_cancellation["is_cancelled"]
    ↪== 1]["total_roll_advert"],
           Roll_Advert_cancellation[Roll_Advert_cancellation["is_cancelled"]
    ↪== 1]["is_cancelled"])
#-----
plt.scatter(Roll_Advert_cancellation[Roll_Advert_cancellation["is_cancelled"]
    ↪== 0]["total_roll_advert"],
           Roll_Advert_cancellation[Roll_Advert_cancellation["is_cancelled"]
    ↪== 0]["is_cancelled"])

plt.yticks((0, 1))
plt.title("Comparing Roll_advertisement against Subscription Cancellation",
    ↪fontsize=14)
plt.xlabel("Number of Rolling advertisement in pages", fontsize=14)
plt.ylabel("1 - is Cancelled, 0 - Not churned", fontsize=14)
plt.grid()
```

Visuals: More rolling advertisement obviously have some effect on unsubscribing the streaming service. There is a parallel trend meaning more advertisement more churning.

1.10.8 Page scatterPlot Summary Findings:

- A. Most widely used pages are 'Home', 'ThumbsUp', 'AddtoPlaylist', 'LogOut', 'RollAdvert', 'AddFriends'.
- B. It seems users who are adding more friends tends to remain with the service.
- C. Adding more song to playlist keeps people with the service.
- D. Rolling advertisement definitely has caused some users to churn out of the service.
- E. ThumbsDown does not influence users decision to leave the service more.
- F. Adding more thumbsUp keeping users with the service almost definitely.
- G. Page error and help needed were not a big factor in unsubscribing the service.

1.11 Browser and Operating platform of the users:

Browsers and operating platforms are not effecting in user's decision for churning the music service in a dominant way. This is just an overview of how these tools might influence potential customers behavior.

```
[106]: def normalize_data(grouped_data):
        grouped_series = grouped_data.set_index(list(grouped_data.columns[:2]))
        temp = grouped_series.unstack('churn').fillna(0)
        df = pd.DataFrame(((temp - temp.min()) / (temp.max() - temp.min()))
        ↪stack()).reset_index()
```

```

    return df
# source: 5 with reference...

```

PLATFORM USAGE

```

[107]: userAgent_count = spark.sql('''
        SELECT userAgent, COUNT(DISTINCT userId) AS user_counts
        FROM df_local_view
        GROUP BY userAgent
        ORDER BY user_counts DESC
    ''').toPandas()

```

```

[108]: df_cleaned.dropDuplicates(["userId"]).groupby(["userAgent", "userId"]).count().
    ↪ sort('userId', ascending=False).limit(5).toPandas()

```

```

[108]:
           userAgent  userId  count
0  "Mozilla/5.0 (Windows NT 6.1; WOW64) AppleWebK...    99      1
1  "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...    98      1
2  "Mozilla/5.0 (Windows NT 5.1) AppleWebKit/537...    97      1
3  "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...    96      1
4  Mozilla/5.0 (Macintosh; Intel Mac OS X 10.9; r...    95      1

```

```

[109]: #platform_dictionary = {'compatible': 'Windows', 'iPad': 'iPad', 'iPhone': '
    ↪ 'iPhone',
#           'Macintosh': 'Mac', 'Windows NT 5.1': 'Windows', 'Windows NT 6.0': '
    ↪ 'Windows', 'Windows NT 6.1': 'Windows',
#           'Windows NT 6.2': 'Windows', 'Windows NT 6.3': 'Windows', 'X11': '
    ↪ 'Linux'}

```

```

[110]: # convert user systems into four categories: Windows, Apple, Linux, other
df_cleaned_churned = df_cleaned_churned.withColumn('oper_platform',
    when((col("userAgent").contains('compatible')) |
    ↪ (col("userAgent").contains('Windows NT 5.1')) | \
        (col("userAgent").contains('Windows NT 6.0')) |
    ↪ (col("userAgent").contains('Windows NT 6.1')) | \
        (col("userAgent").contains('Windows NT 6.2')) |
    ↪ (col("userAgent").contains('Windows NT 6.3')), 'Windows')\
        .when((col("userAgent").contains('Macintosh')) |
    ↪ (col("userAgent").contains('iPhone'))\
        |
    ↪ (col("userAgent").contains('iPad')), 'Apple')\
        .when((col("userAgent").contains('X11')) |
    ↪ (col("userAgent").contains('Linux')), 'Linux')
        .otherwise('other'))
#df_cleaned_churned.limit(2).toPandas()

```

```
[111]: platform_data = df_cleaned_churned.select(['oper_platform', 'churn']).
        ↳groupby(['oper_platform', 'churn']).\
                                count().sort('oper_platform').toPandas()
platform_data.head(3)
```

```
[111]:   oper_platform  churn  count
0         Apple      1   18365
1         Apple      0  106957
2         Linux      1    1762
```

BROWSERS USAGE


```
[112]: # convert user systems into four categories: Windows, Apple, Linux, other
df_cleaned_churned = df_cleaned_churned.withColumn('browsers',
        when(col("userAgent").contains('Chrome'), 'Chrome')
        .when((col("userAgent").contains('Firefox')), 'Firefox')
        .when(col("userAgent").contains('Safari'), 'Safari')
        .otherwise('IE'))

#df_cleaned_churned.limit(2).toPandas()
```

```
[113]: browser_selection = df_cleaned_churned.dropDuplicates(["userId"]).
        ↳groupby(["userId", "browsers"]).count().\
                                                    ↳
        ↳sort('userId', ascending=False).toPandas()
fig, ax = plt.subplots(figsize=(14,5))

ax = browser_selection.groupby('browsers')['count'].sum().
        ↳sort_values(ascending=False).plot(kind='bar',\
        ↳color=(['salmon', 'olivedrab', 'tan', 'plum']))
ax.set_facecolor('ivory')

ax.set_xlabel('Browsers', fontsize=13)
ax.set_ylabel('Browser count', fontsize=13)
ax.set_title("Browser choices among users", fontsize=14)
ax.set_xticklabels(ax.get_xticklabels(), rotation=0)
ax.grid()
```



```
[114]: browsers_data = df_cleaned_churned.select(['browsers', 'churn']).
      ↪groupby(['browsers', 'churn']).\
      count().sort('browsers').toPandas()
browsers_data.head(3)
```

```
[114]:   browsers  churn  count
0   Chrome      1  20561
1   Chrome      0 125030
2  Firefox      1  14847
```

```
[115]: platform_data = normalize_data(platform_data)
browsers_data = normalize_data(browsers_data)

plt.figure(figsize=(20,5))
# -----Browsers-----
plt.subplot(131)
sns.barplot( x="browsers", y="count", hue="churn", data=browsers_data);

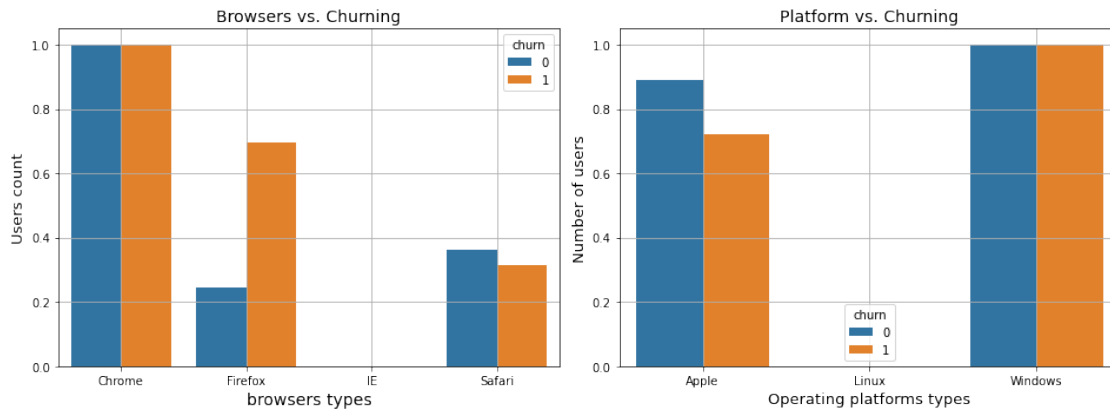
plt.title("Browsers vs. Churning", fontsize=14)
plt.xlabel("browsers types", fontsize=14)
plt.ylabel("Users count", fontsize=13)
plt.grid()
ax.set_facecolor('ivory')

#-----Platform-----
plt.subplot(132)
sns.barplot(x='oper_platform', y='count', hue='churn', data=platform_data)

plt.title('Platform vs. Churning', fontsize=14)
plt.xlabel("Operating platforms types", fontsize=13)
plt.ylabel("Number of users", fontsize=13)
```

```
plt.grid()
ax.set_facecolor('ivory')

plt.tight_layout()
```



Browser and Platform distribution: It is obvious that there are more Chrome, Firefox than Safari browsers. Windows operating system is more widely used than Apple systems thereby churning is correlated just by the market share nothing more.

Scaled summary visual: Typically 'Google' has a dominant market share in browsers market and it is obvious that more people have churned out from Chrome based devices than Apple-Safari or Firefox devices. There was no drastic relations of churning with one kind of browsers or the others. This visual offers a view where data is scaled.

1.12 Churning effect on hourly, daily, monthly basis among users:

```
[116]: get_month = udf(lambda x: datetime.datetime.fromtimestamp(x / 1000.0).month,
    ↪IntegerType())
get_day = udf(lambda x: datetime.datetime.fromtimestamp(x / 1000.0).day,
    ↪IntegerType())
get_hour = udf(lambda x: datetime.datetime.fromtimestamp(x / 1000.0).hour,
    ↪IntegerType())

get_weekday = udf(lambda x: datetime.datetime.fromtimestamp(x / 1000.0).
    ↪strftime('%w'))
```

```
[117]: df_cleaned_churned = df_cleaned_churned.withColumn("month",
    ↪get_month(df_cleaned_churned.ts))
df_cleaned_churned = df_cleaned_churned.withColumn("day",
    ↪get_day(df_cleaned_churned.ts))
df_cleaned_churned = df_cleaned_churned.withColumn("hour",
    ↪get_hour(df_cleaned_churned.ts))
```

```
df_cleaned_churned = df_cleaned_churned.withColumn("weekday",   
↳get_weekday(df_cleaned_churned.ts))
```

```
[118]: df_cleaned_churned.take(1)
```

```
[118]: [Row(artist='Sleeping With Sirens', auth='Logged In', firstName='Darianna',  
gender='F', itemInSession=0, lastName='Carpenter', length=202.97098,  
level='free', location='Bridgeport-Stamford-Norwalk, CT', method='PUT',  
page='NextSong', registration=1538016340000, sessionId=31, song='Captain Tyin  
Knots VS Mr Walkway (No Way)', status=200, ts=1539003534000,  
userAgent='"Mozilla/5.0 (iPhone; CPU iPhone OS 7_1_2 like Mac OS X)  
AppleWebKit/537.51.2 (KHTML, like Gecko) Version/7.0 Mobile/11D257  
Safari/9537.53"', userId='100010', churn=0, oper_platform='Apple',  
browsers='Safari', month=10, day=8, hour=5, weekday='1')]
```


```
[119]: print("Statistical analysis of month column: ")  
df_cleaned_churned.select("month").describe().show()
```

Statistical analysis of month column:

```
+-----+-----+  
|summary|          month|  
+-----+-----+  
|  count|          278154|  
|   mean|10.472540391294032|  
| stddev|0.5024405873732121|  
|    min|              9|  
|    max|             12|  
+-----+-----+
```

```
[120]: month_data = df_cleaned_churned.select(["churn", "month"]).groupby(["churn",  
↳"month"]).count().sort("month").toPandas()  
month_data = normalize_data(month_data)  
  
fig, ax = plt.subplots(figsize=(14, 5), edgecolor='k', facecolor='wheat')  
  
sns.barplot( x="month",y="count", hue="churn", data=month_data )  
ax.set_title("Churning and remained users from October to December months",  
↳fontSize=14)
```

```
[120]: Text(0.5, 1.0, 'Churning and remained users from October to December months')
```




Monthly churning: Monthly data available only for 3 months which is inadequate to decide a trend of churning behavior in rest of the years.

```
[121]: day_data = df_cleaned_churned.select(["churn", "day"]).groupby(["churn", "day"]).count().sort("day").toPandas()
day_data = normalize_data(day_data)

fig, ax = plt.subplots(figsize=(14, 5), edgecolor='k', facecolor='wheat')

ax = sns.barplot( x="day", y="count", hue="churn", data=day_data )
ax.set_xlabel("Day 1 to 31", fontsize=14)
ax.set_ylabel("Day count", fontsize=13)
ax.set_title("Churned and Remained users on Daily basis", fontsize=14)
```

[121]: Text(0.5, 1.0, 'Churned and Remained users on Daily basis')




- An interesting observation is there are more churning is happening until the 15th of each month. From 16th on to the end of the month churning reduce down to half of retained/stayed customers.

```
[122]: hour_data = df_cleaned_churned.select(["churn", "hour"]).groupby(["churn", "hour"]).count().sort("hour").toPandas()
hour_data = normalize_data(hour_data)

fig, ax = plt.subplots(figsize=(14, 5), edgecolor='k', facecolor='moccasin')
ax = sns.barplot(x="hour", y="count", hue="churn", data=hour_data);
ax.set_xlabel("24 hours 0-23", fontsize=14)
ax.set_ylabel("Hour Count", fontsize=14)
ax.set_title("Churned and Remained users on Hourly basis", fontsize=14)
```


```
[122]: Text(0.5, 1.0, 'Churned and Remained users on Hourly basis')
```



- Most number of churning happen at the 10 am 11 am morning time then slowly it goes down and again 16th and 17th hour goes up again.

```
[123]: weekday_data = df_cleaned_churned.select(["churn", "weekday"]).groupby(["churn", "weekday"]).count().sort("weekday").toPandas()
weekday_data = normalize_data(weekday_data)

fig, ax = plt.subplots(figsize=(14, 5), edgecolor='k', facecolor='moccasin')
ax = sns.barplot(x="weekday", y="count", hue="churn", data=weekday_data);
ax.set_xlabel("Weekdays Monday to Sunday", fontsize=14)
ax.set_ylabel("Weekdays count", fontsize=14)
ax.set_title("Churned and Remained users on Weekday basis", fontsize=14)
ax.set_xticklabels(['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun'], rotation=0);
```

Day wise churning: It seems Monday have least or no churning next is Sunday. The rest of the weekdays stays almost the same rate of churning transition sequence.

```
[124]: from pyspark.sql.functions import split
df_cleaned_churned = df_cleaned_churned.withColumn("city",
↳split(col("location"), ",").getItem(0)).\
withColumn("state", split(col("location"), ",").
↳getItem(1))
```

```
[125]: # source: https://www.datasciencemadesimple.com/
↳string-split-of-the-columns-in-pyspark/
```

```
[126]: city_data = df_cleaned_churned.dropDuplicates(["userId"]).
↳groupby(["city", "churn"]).count().\

↳sort("city", ascending=True).toPandas()
#city_data = normalize_data(city_data)
```

```
[127]: state_data = df_cleaned_churned.dropDuplicates(["userId"]).
↳groupby(["state", "churn"]).count().\


↳sort("state", ascending=True).toPandas()
#state_data = normalize_data(state_data)
```

```
[128]: plt.figure(figsize=(20,10))

# State wise churning
plt.subplot(131)
plt.title('State 0-remained 1-churned', fontsize=14)
sns.barplot( x="count", y="state", hue="churn", data=state_data[:55]);
```

```
# city wise churning
plt.subplot(132)
plt.title('City 0-remained 1-churned', fontsize=14)
sns.barplot( x="count",y="city", hue="churn", data=city_data[:55]);

plt.tight_layout()
plt.show()
```



Location and churning: This is very a clean visualization of users' churning trend based on city and state.

2 Feature Engineering processes:

Now that we have a better grasp of the data that we're given, lets start thinking about feature engineering and manipulation. We know feature engineering involves the application of business knowledge, statistics to transform data into a format that can be directly used by machine learning models. Irrespective of the algorithm used, feature engineering drives model performances to generate meaningful insights, and ultimately critical solve business problems. Primary objectives are..

- Manipulating existing features with analytical relevancy to other features.
- Redesigning new and complex but needed feature relationship.
- Finding complex inter-relational important values among existing features.

We will do these options by aggregating features and joining multiple tables into a single dataframe while using statistical transformations with relevant relational operations. On these progression we will convert all object columns into numerical ones.

The ultimate leverage of feature manipulation is to predict customers who will likely churn out of the service with the highest probability. If we can correctly identify those customers, we might be able to retain them as a continued customer with new promotions.

```
[271]: df_cleaned_churned.printSchema()
```

```
root
|-- artist: string (nullable = true)
|-- auth: string (nullable = true)
|-- firstName: string (nullable = true)
|-- gender: string (nullable = true)
|-- itemInSession: long (nullable = true)
|-- lastName: string (nullable = true)
|-- length: double (nullable = true)
|-- level: string (nullable = true)
|-- location: string (nullable = true)
|-- method: string (nullable = true)
|-- page: string (nullable = true)
|-- registration: long (nullable = true)
|-- sessionId: long (nullable = true)
|-- song: string (nullable = true)
|-- status: long (nullable = true)
|-- ts: long (nullable = true)
|-- userAgent: string (nullable = true)
|-- userId: string (nullable = true)
|-- churn: long (nullable = true)
|-- oper_platform: string (nullable = false)
|-- browsers: string (nullable = false)
|-- month: integer (nullable = true)
|-- day: integer (nullable = true)
|-- hour: integer (nullable = true)
|-- weekday: string (nullable = true)
|-- city: string (nullable = true)
|-- state: string (nullable = true)
```

Quick summary of feature transformations:

- In the first function we renamed existing column and added 4 new columns
- In the second function we renamed and recreated 7 new columns
- Aggregating features as relevant and joining those two functions
- Renaming multiple of the sub-selected columns

- Then we applied log-transformation on large numeric datatypes
- At this point the dataset is ready for Vectorization and Scaling processes

Note: Computation time with PySpark was a factor for not trying more new feature engineering aspects.

2.0.1 Total number of columns we'd be working with

```
[140]: features_chosen = df_cleaned_churned.  
        ↪select(['userId', 'auth', 'gender', 'level', 'page', 'ts', 'registration', 'length', 'hour', 'day'])  
        features_chosen.show(2)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+-+
|userId|    auth|gender|level|    page|          ts| registration|
length|hour|day|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+-+
|    30|Logged In|    M| paid|NextSong|1538352117000|1538173362000|277.89016|
17| 30|
|    9|Logged In|    M| free|NextSong|1538352180000|1538331630000|236.09424|
17| 30|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+-+
only showing top 2 rows
```

```
[141]: def get_basic_features():  
        '''  
        INPUT:  
        There will be no input  
        We will be sub-selecting 7 relevent columns  
        We will be renaming and recreating new columns....  
  
        OUTPUT:  
        Final output will contain 'userId' in correlation with newly formatted_  
        ↪columns  
        '''  
        basic_features = df_cleaned_churned.  
        ↪select(['userId', 'auth', 'gender', 'level', 'ts', 'registration', 'length', 'hour', 'day'])  
        #-----  
        male_gender = udf(lambda x: 1 if x=='M' else 0, IntegerType())  
  
        # add downgrade variable  
        basic_features = basic_features.withColumn("Male_Female",  
        ↪male_gender("gender"))
```

```

#-----
auth_selection = udf(lambda x: 1 if x=='Logged In' else 0, IntegerType())
basic_features = basic_features.withColumn("LogIn_Cancelled",
→auth_selection("auth"))

#-----
# define function
level_event = udf(lambda x: 1 if x == "paid" else 0, IntegerType())

# add gender variable
basic_features = basic_features.withColumn("paid_free",
→level_event("level"))

#-----
calculate_time = udf(lambda x,y: x-y if x>=y else 0, LongType())

# add gender variable
result = basic_features.withColumn("timeSinceRegistration",
→calculate_time("ts", "registration"))

return result

```

```

[142]: main_features_only = get_basic_features()
      #basic_features_only.limit(3).toPandas()

```

```

[143]: main_features_only.dtypes

```

```

[143]: [('userId', 'string'),
      ('auth', 'string'),
      ('gender', 'string'),
      ('level', 'string'),
      ('ts', 'bigint'),
      ('registration', 'bigint'),
      ('length', 'double'),
      ('hour', 'int'),
      ('day', 'int'),
      ('Male_Female', 'int'),
      ('LogIn_Cancelled', 'int'),
      ('paid_free', 'int'),
      ('timeSinceRegistration', 'bigint')]

```

```

[144]: # deleting 'string' features to keep only numerical value features
      columns = ['auth', 'gender', 'level']
      main_features_only = main_features_only.drop(*columns)
      main_features_only.show(3)

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+

```

```

--+-----+-----+
|userId|          ts| registration|
length|hour|day|Male_Female|LogIn_Cancelled|paid_free|timeSinceRegistration|
+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+
|    30|1538352117000|1538173362000|277.89016|  17| 30|          1|
1|          1|          178755000|
|    9|1538352180000|1538331630000|236.09424|  17| 30|          1|
1|          0|          20550000|
|    30|1538352394000|1538173362000| 282.8273|  17| 30|          1|
1|          1|          179032000|
+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+
only showing top 3 rows

```

2.0.2 Page features reengineering

```
[145]: df_cleaned_churned.select('page').dropDuplicates().collect()
```

```
[145]: [Row(page='Cancel'),
Row(page='Submit Downgrade'),
Row(page='Thumbs Down'),
Row(page='Home'),
Row(page='Downgrade'),
Row(page='Roll Advert'),
Row(page='Logout'),
Row(page='Save Settings'),
Row(page='Cancellation Confirmation'),
Row(page='About'),
Row(page='Settings'),
Row(page='Add to Playlist'),
Row(page='Add Friend'),
Row(page='NextSong'),
Row(page='Thumbs Up'),
Row(page='Help'),
Row(page='Upgrade'),
Row(page='Error'),
Row(page='Submit Upgrade')]
```

```
[146]: def get_page_features():
'''
INPUT:
Sub-selecting 3 relevent features
We will be manipulating page columns into multiple sub category columns
to make sure all new columns are in number type dataset
'''
```

OUTPUT:

Final output will contain 'userId' and newly formatted feature columns
'''

```
page_features = df_cleaned_churned.select(['userId', 'page', 'sessionId',  
↳ 'artist'])  
  
# add downgrade variable  
downgraded_event = udf(lambda x: 1 if x == "Submit Downgrade" else 0,  
↳ IntegerType())  
  
page_features = page_features.withColumn("downgraded",  
↳ downgraded_event("page"))  
  
#-----Roll  
↳ Advert-----  
roll_advert_event = udf(lambda x: 1 if x == "Roll Advert" else 0,  
↳ IntegerType())  
  
page_features = page_features.withColumn("rolling_Ad",  
↳ roll_advert_event("page"))  
  
# # define  
↳ Playlist-----  
playlist_event = udf(lambda x: 1 if x == "Add to Playlist" else 0,  
↳ IntegerType())  
  
# add playlist variable  
page_features = page_features.withColumn("playlistSongs",  
↳ playlist_event("page"))  
  
#-----Add  
↳ friends-----  
Add_Friend_event = udf(lambda x: 1 if x == "Add Friend" else 0,  
↳ IntegerType())  
page_features = page_features.withColumn('addedFreind',  
↳ Add_Friend_event('page'))  
  
# add thumbs-up variable-----  
thumbs_up_event = udf(lambda x: 1 if x == "Thumbs Up" else 0, IntegerType())  
page_features = page_features.withColumn('thumbsUp',  
↳ thumbs_up_event('page'))  
  
# define  
↳ function-----
```

```

thumbs_down_event = udf(lambda x: 1 if x == "Thumbs Down" else 0,
↳IntegerType())
page_features = page_features.withColumn('thumbsDown',
↳thumbs_down_event('page'))

↳
↳#=====
# define function-----
song_event = udf(lambda x: 1 if x == "NextSong" else 0, IntegerType())

# add songs variable
page_features = page_features.withColumn("songChoice", song_event("page"))

###-----PAGE --Interrelated↳
↳sessions-----##
num_total_song = page_features.filter(page_features.page=='NextSong').
↳select(["userId", "songChoice"]).\
dropDuplicates(["userId"]).
↳groupby(["userId"]).count()

num_total_song = num_total_song.selectExpr("userId as userId", "count as
↳num_total_song")

↳
↳#-----
↳
# Number of session by each users
num_sessions = page_features.filter(page_features.page == 'NextSong').
↳select(['userId', 'sessionId']).\
↳dropDuplicates(["userId", "sessionId"]).groupby(["userId"]).count()

num_sessions = num_sessions.selectExpr("userId as userId", "count as
↳user_sessions")

↳
↳#-----
#unique artists artist
num_artist = page_features.filter(page_features.page=='NextSong').
↳select(["userId", "artist"]).\
↳dropDuplicates(["userId", "artist"]).groupby(["userId"]).count()

num_artist = num_artist.selectExpr("userId as userId", "count as num_artist")

```



```

    combined_features = page_features.join(num_total_song, on='userId').\
                                join(num_sessions, on='userId').
    join(num_artist, on='userId')
    # return a combined all modified columns
    return combined_features

```

```

[147]: page_features_subset = get_page_features()
columns = ['page', 'artist']
page_features_subset = page_features_subset.drop(*columns)
page_features_subset.show(3)

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+
|userId|sessionId|downgraded|rolling_Ad|playlistSongs|addedFreind|thumbsUp|thumb
sDown|songChoice|num_total_song|user_sessions|num_artist|
+-----+-----+-----+-----+-----+-----+-----+-----+
|100010|      31|      0|      0|      0|      0|      0|
0|      1|      1|      7|     252|
|100010|      31|      0|      0|      0|      0|      0|
0|      1|      1|      7|     252|
|100010|      31|      0|      0|      0|      0|      0|
0|      1|      1|      7|     252|
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 3 rows

```

```

[148]: page_features_subset.dtypes

```

```

[148]: [('userId', 'string'),
        ('sessionId', 'bigint'),
        ('downgraded', 'int'),
        ('rolling_Ad', 'int'),
        ('playlistSongs', 'int'),
        ('addedFreind', 'int'),
        ('thumbsUp', 'int'),
        ('thumbsDown', 'int'),
        ('songChoice', 'int'),
        ('num_total_song', 'bigint'),
        ('user_sessions', 'bigint'),
        ('num_artist', 'bigint')]

```

```

[149]: page_features_subset.limit(3).toPandas()

```

```
[149]:   userId  sessionId  downgraded  rolling_Ad  playlistSongs  addedFreind  \
0  100010         31           0           0           0           0
1  100010         31           0           0           0           0
2  100010         31           0           0           0           0

      thumbsUp  thumbsDown  songChoice  num_total_song  user_sessions  num_artist
0           0           0           1           1           7           252
1           0           0           1           1           7           252
2           0           0           1           1           7           252
```

2.0.3 Combining all the data subsets

```
[150]: userId_churn = df_cleaned_churned.select(['userId', 'churn']).dropDuplicates()
main_page_combined = main_features_only.join(page_features_subset,
↳ on='userId', how='inner').\
                                                    join(userId_churn, on='userId',
↳ how='inner')
```

```
[151]: main_page_combined.show(3)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
---+-----+
|userId|          ts| registration|   length|hour|day|Male_Female|LogIn_Cancelled|paid_free|timeSinceRegistration|sessionId|downgraded|rolling_Ad|playlistSongs|addedFreind|thumbsUp|thumbsDown|songChoice|num_total_song|user_sessions|num_artist|churn|
+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
---+-----+
|100010|1539003534000|1538016340000|202.97098|  5| 8|          0|
1|          0|          987194000|          31|          0|          0|          0|
0|          0|          0|          1|          1|          7|          252|          0|
|100010|1539003534000|1538016340000|202.97098|  5| 8|          0|
1|          0|          987194000|          31|          0|          0|          0|
0|          0|          0|          1|          1|          7|          252|          0|
|100010|1539003534000|1538016340000|202.97098|  5| 8|          0|
1|          0|          987194000|          31|          0|          0|          0|
0|          0|          0|          1|          1|          7|          252|          0|
+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
---+-----+
only showing top 3 rows
```

```
[152]: main_page_combined.dtypes
```

```
[152]: [('userId', 'string'),
        ('ts', 'bigint'),
        ('registration', 'bigint'),
        ('length', 'double'),
        ('hour', 'int'),
        ('day', 'int'),
        ('Male_Female', 'int'),
        ('LogIn_Cancelled', 'int'),
        ('paid_free', 'int'),
        ('timeSinceRegistration', 'bigint'),
        ('sessionId', 'bigint'),
        ('downgraded', 'int'),
        ('rolling_Ad', 'int'),
        ('playlistSongs', 'int'),
        ('addedFreind', 'int'),
        ('thumbsUp', 'int'),
        ('thumbsDown', 'int'),
        ('songChoice', 'int'),
        ('num_total_song', 'bigint'),
        ('user_sessions', 'bigint'),
        ('num_artist', 'bigint'),
        ('churn', 'bigint')]
```

2.0.4 Temporal aggregation and renaming features variables for feature minimization

```
[153]: # Feature aggregation columns by userId
all_features_combined = main_page_combined.groupby(main_page_combined.userId).\
    agg({"downgraded": "max",
        "Male_Female": "max",
        "paid_free": "max",
        "hour": "sum",
        "day": "sum",

        "timeSinceRegistration": "max",
        "rolling_Ad": "sum",
        "playlistSongs": "sum",
        "thumbsUp": "sum",
        "thumbsDown": "sum",
        "length": "avg",

        "songChoice": "sum",
```

```

        "num_total_song": "sum",
        "num_artist": "sum",
        "user_sessions": "sum",
        "churn": "max"}}

```

RENAMING AGGREGATED FEATURES

```

[154]: # Source: https://amiradata.com/pyspark-rename-column-on-pyspark-dataframe/
# renaming aggregated feture columns
users_features_inclusive = all_features_combined.select(col("userId"),
    col("max(downgraded)").alias("downgraded"),
    col("max(Male_Female)").alias("Male_Female"),
    col("max(paid_free)").alias("paid_free"),
    col("sum(hour)").alias("total_hour"),
    col("sum(day)").alias("num_days"),

    col("max(timeSinceRegistration)").
↪alias("timeSinceRegistration"),
    col("sum(rolling_Ad)").alias("rolling_Ad"),
    col("sum(playlistSongs)").
↪alias("playlistSongs"),
    col("sum(thumbsUp)").alias("thumbsUp"),

    col("sum(thumbsDown)").alias("thumbsDown"),
    col("avg(length)").alias("length"),
    col("sum(songChoice)").alias("songChoice"),
    col("sum(num_total_song)").
↪alias("num_total_song"),
    col("sum(num_artist)").alias("num_artist"),
    col("sum(user_sessions)").
↪alias("user_sessions"),
    col("max(churn)").alias("churn"))

# present result
users_features_inclusive.show(3)

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+
|userId|downgraded|Male_Female|paid_free|total_hour|num_days|timeSinceRegistrati
on|rolling_Ad|playlistSongs|thumbsUp|thumbsDown|
length|songChoice|num_total_song|num_artist|user_sessions|churn|
+-----+-----+-----+-----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+
|100010|          0|          0|          0|  1497711| 1711071|

```

```

4807612000|      19812|      2667|      6477|      1905|243.42144490910485|
104775|      145161| 36580572|      1016127|      0|
|200002|      0|      1|      1| 2120676| 3716634|
6054448000|      3318|      3792|      9954|      2844|242.91699209303025|
183438|      224676| 76165164|      1348056|      0|
| 125|      0|      1|      0|      2541|      1331|
6161779000|      11|      0|      0|      0|261.13913749999995|
88|      121|      968|      121|      1|
+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+
only showing top 3 rows

```

```
[155]: users_features_inclusive.dtypes
```

```
[155]: [('userId', 'string'),
        ('downgraded', 'int'),
        ('Male_Female', 'int'),
        ('paid_free', 'int'),
        ('total_hour', 'bigint'),
        ('num_days', 'bigint'),
        ('timeSinceRegistration', 'bigint'),
        ('rolling_Ad', 'bigint'),
        ('playlistSongs', 'bigint'),
        ('thumbsUp', 'bigint'),
        ('thumbsDown', 'bigint'),
        ('length', 'double'),
        ('songChoice', 'bigint'),
        ('num_total_song', 'bigint'),
        ('num_artist', 'bigint'),
        ('user_sessions', 'bigint'),
        ('churn', 'bigint')]
```

```
[156]: #users_features_inclusive.limit(5).toPandas()
```

```
[157]: features_in_pandas = users_features_inclusive.toPandas()
features_in_pandas.groupby('churn').mean()
```

```
[157]:
```

	downgraded	Male_Female	paid_free	total_hour	num_days	\
churn						
0	0.231214	0.514451	0.745665	4.323380e+07	6.174784e+07	
1	0.173077	0.615385	0.692308	1.679000e+07	1.910449e+07	

	timeSinceRegistration	rolling_Ad	playlistSongs	thumbsUp	\
churn					
0	7.484022e+09	36143.358382	91853.242775	177210.780347	

```

1          4.951238e+09  20262.923077  34437.269231  62874.634615

      thumbsDown      length  songChoice  num_total_song  num_artist \
churn
0      33293.075145  249.141254  3.170549e+06    3.838546e+06  7.721456e+09
1      14020.269231  248.307865  1.200802e+06    1.456754e+06  1.815173e+09

      user_sessions
churn
0      1.827633e+08
1      3.069667e+07

```

2.0.5 Applying log transformation on large-value columns

```

[158]: # Check average feature value for each target value
features_in_pandas = users_features_inclusive.toPandas()
features_in_pandas.groupby('churn').mean()

```

```

[158]:      downgraded  Male_Female  paid_free  total_hour  num_days \
churn
0      0.231214      0.514451   0.745665  4.323380e+07  6.174784e+07
1      0.173077      0.615385   0.692308  1.679000e+07  1.910449e+07

      timeSinceRegistration  rolling_Ad  playlistSongs  thumbsUp \
churn
0      7.484022e+09  36143.358382   91853.242775  177210.780347
1      4.951238e+09  20262.923077   34437.269231   62874.634615

      thumbsDown      length  songChoice  num_total_song  num_artist \
churn
0      33293.075145  249.141254  3.170549e+06    3.838546e+06  7.721456e+09
1      14020.269231  248.307865  1.200802e+06    1.456754e+06  1.815173e+09

      user_sessions
churn
0      1.827633e+08
1      3.069667e+07

```

```

[159]: # apply log to all continuous features
features_streamlined = users_features_inclusive.select(col("Male_Female"),
                                                    col("downgraded"), col("paid_free"),
                                                    F.log(col("timeSinceRegistration") + 1).
→alias("timeSinceRegistration"),
                                                    F.log(col("playlistSongs") + 1).alias("playlistSongs"),
                                                    F.log(col("rolling_Ad") + 1).alias("rolling_Add"),
                                                    F.log(col("thumbsUp") + 1).alias("thumbsUp"),

```

```

F.log(col("thumbsDown") + 1).
↳alias("thumbsDown"),
F.log(col("length") + 1).alias("length"),
F.log(col("songChoice") + 1).
↳alias("songChoice"),
F.log(col("num_total_song") + 1).
↳alias("num_total_song"),
F.log(col("num_artist") + 1).
↳alias("num_artist"),
F.log(col("user_sessions") + 1).
↳alias("user_sessions"),
F.log(col("num_days") + 1).
↳alias("num_days"),
F.log(col("total_hour") + 1).
↳alias("total_hour"),
col("churn"))

```

```
[160]: features_streamlined.dtypes
```

```

[160]: [('Male_Female', 'int'),
        ('downgraded', 'int'),
        ('paid_free', 'int'),
        ('timeSinceRegistration', 'double'),
        ('playlistSongs', 'double'),
        ('rolling_Add', 'double'),
        ('thumbsUp', 'double'),
        ('thumbsDown', 'double'),
        ('length', 'double'),
        ('songChoice', 'double'),
        ('num_total_song', 'double'),
        ('num_artist', 'double'),
        ('user_sessions', 'double'),
        ('num_days', 'double'),
        ('total_hour', 'double'),
        ('churn', 'bigint')]

```

```

[161]: # to view how newly created features affect on churning..
features_streamlined.toPandas().groupby('churn').mean().head(3)

```

```

[161]:      Male_Female  downgraded  paid_free  timeSinceRegistration \
churn
0          0.514451    0.231214    0.745665             22.632942
1          0.615385    0.173077    0.692308             22.103498

      playlistSongs  rolling_Add  thumbsUp  thumbsDown  length \
churn

```

0	9.221374	8.168837	10.007560	7.999916	5.521801
1	8.369453	8.093355	8.941163	7.557110	5.518217

	songChoice	num_total_song	num_artist	user_sessions	num_days \
churn					
0	12.966921	13.187744	19.321687	15.461751	15.929829
1	11.995807	12.263976	17.950117	14.191401	14.636405

	total_hour
churn	
0	15.583733
1	14.675564

```
[162]: # display size after all feature transformation
print('We got ', users_features_inclusive.count(), 'rows and ',
      len(users_features_inclusive.columns), 'columns')
```

We got 225 rows and 17 columns

```
[163]: # separating the columns
final_dataSet_columns = features_streamlined.columns
print("Before removing target columns: ['churn'] ")
final_dataSet_columns
```

Before removing target columns: ['churn']

```
[163]: ['Male_Female',
        'downgraded',
        'paid_free',
        'timeSinceRegistration',
        'playlistSongs',
        'rolling_Add',
        'thumbsUp',
        'thumbsDown',
        'length',
        'songChoice',
        'num_total_song',
        'num_artist',
        'user_sessions',
        'num_days',
        'total_hour',
        'churn']
```

```
[164]: # create pandas version of dataset after log-transformation
#features_streamlined_pandas = features_streamlined.toPandas()
#features_streamlined_pandas.groupby('churn').mean()
```


2.0.6 Applying 'Vectorization + Scaling' on the final dataset

Note: Some algorithms like Logistics Regression and SVC performs poorly if the dataset is not scaled. So we will do scaling of the vectorized dataset to be in safer side.

```
[169]: # removing the target columns
final_dataSet_columns.remove('churn')
print("After removing target columns:['churn']")
final_dataSet_columns
```

After removing target columns:['churn']

```
[169]: ['Male_Female',
        'downgraded',
        'paid_free',
        'timeSinceRegistration',
        'playlistSongs',
        'rolling_Add',
        'thumbsUp',
        'thumbsDown',
        'length',
        'songChoice',
        'num_total_song',
        'num_artist',
        'user_sessions',
        'num_days',
        'total_hour']
```

```
[170]: # Applying vectorization on the redefined dataset
assembler = VectorAssembler(inputCols = final_dataSet_columns,
    ↪outputCol="Num_Features")
data = assembler.setHandleInvalid('skip').transform(features_streamlined)
```

```
[171]: # Scaling to mean 0 and unit std dev
scaler = StandardScaler(inputCol='Num_Features', outputCol='features_scaled',
    ↪withMean=True, withStd=True)
ScaledDataSet = scaler.fit(data)
data = ScaledDataSet.transform(data)
```

```
[172]: data.limit(3).toPandas()
```

```
[172]:   Male_Female  downgraded  paid_free  timeSinceRegistration  playlistSongs \
0           0           0           0           22.293466           7.889084
1           1           0           1           22.524059           8.240913
2           1           0           0           22.541631           0.000000

   rolling_Add  thumbsUp  thumbsDown   length  songChoice  num_total_song \
```

0	9.894094	8.776167	7.552762	5.498894	11.559580	11.885606
1	8.107419	9.205830	7.953318	5.496828	12.119637	12.322419
2	2.484907	0.000000	0.000000	5.568875	4.488636	4.804021

	num_artist	user_sessions	num_days	total_hour	churn	\
0	17.415028	13.831510	14.352631	14.219449	0	
1	18.148415	14.114175	15.128329	14.567246	0	
2	6.876265	4.804021	7.194437	7.840706	1	

	Num_Features	\
0	[0.0, 0.0, 0.0, 22.293466332295722, 7.88908440...	
1	[1.0, 0.0, 1.0, 22.524059045633912, 8.24091254...	
2	[1.0, 0.0, 0.0, 22.54163137166111, 0.0, 2.4849...	

	features_scaled
0	[-1.0762390990345028, -0.5264710031632485, -1...
1	[0.9250319528891594, -0.5264710031632485, 0.60...
2	[0.9250319528891594, -0.5264710031632485, -1.6...

```
[173]: # renaming target and features-columns
final_dataSet = data.select(data.churn.alias("label"), data.features_scaled.
    ↪alias("features"))
final_dataSet.show(3)
```

```
+-----+-----+
|label|          features|
+-----+-----+
|    0|[-1.0762390990345...|
|    0|[0.92503195288915...|
|    1|[0.92503195288915...|
+-----+-----+
only showing top 3 rows
```

```
[174]: # displaying label count
final_dataSet.groupby(final_dataSet.label).count().show()
```

```
+-----+-----+
|label|count|
+-----+-----+
|    0|   173|
|    1|    52|
+-----+-----+
```

2.0.7 Balancing minority (churning) class

- Since we have a highly imbalanced target value churned 52 which is around 1/3 of the non-churned 173, we will apply undersampling in pursuasion of creating a balanced dataset.

```
[175]: def undersample(df, minority, majority):  
    '''  
    Implement undersample on dataset, return a balanced dataset.  
    '''  
    # size of minority class(0)  
    minoritySize = df.where(df.label == minority).count()  
  
    # two classes with the same size  
    df_minority = df.where(df.label == minority)  
    df_majority = df.where(df.label == majority).sample(1.0, seed=7).  
    ↪limit(minoritySize)  
  
    # concatenate them together  
    result = df_minority.union(df_majority)  
  
    #shuffle data  
    result = result.orderBy(rand())  
    return result
```

```
[176]: # applying balancing function  
balanced_dataSet = undersample(final_dataSet, 1, 0)  
  
# finding the lable count  
balanced_dataSet.groupby(balanced_dataSet.label).count().show()
```

```
+-----+-----+  
|label|count|  
+-----+-----+  
|    0|   52|  
|    1|   52|  
+-----+-----+
```

```
[177]: balanced_dataSet.toPandas().head(4)
```

```
[177]:   label  features  
0      0  [0.9250319528891594, -0.5264710031632485, -1.6...  
1      0  [-1.0762390990345028, -0.5264710031632485, -1...  
2      0  [0.9250319528891594, -0.5264710031632485, 0.60...  
3      0  [0.9250319528891594, -0.5264710031632485, -1.6...
```

2.1 ML modeling with DataSet splitting:

```
[189]: df_train, df_test = balanced_dataSet.randomSplit([0.8, 0.2], seed=15)
```

```
[190]: print("Now the shape of df_train dataset: ",(df_train.count(), len(df_train.
    ↪columns)))
```

Now the shape of df_train dataset: (87, 2)

```
[191]: print("Now the shape of df_test dataset: ",(df_test.count(), len(df_test.
    ↪columns)))
```

Now the shape of df_test dataset: (17, 2)

Since our target variable is a binary classification choice, I've decided to use accuracy and f-1 score. Considering we have a balanced dataset makes accuracy is the prime metric to observe also they're easy to interpret.

- Accuracy describes how often our model is correct regardless of the type of errors it makes.
- F-1 score balances the tradeoff between precision and recall.
- precision (how often is the model correct over every "positive" prediction) and
- recall (how many of the total "positive" instances were identified correctly).
- Confusion Matrix to visualize the performance of an algorithm.

```
[192]: import time
evaluator= MulticlassClassificationEvaluator(predictionCol="prediction")
```

```
[193]: # securing metrics to choose
accuracy_evaluator = MulticlassClassificationEvaluator(metricName='accuracy')
f1_evaluator = MulticlassClassificationEvaluator(metricName='f1')
```

```
[194]: # Initialize five models
clf_LR = LogisticRegression(maxIter=40)
clf_SVC = LinearSVC()
clf_RF = RandomForestClassifier(numTrees = 10, seed=35)
clf_DT = DecisionTreeClassifier(seed=30)
```

```
[195]: def model_evalutaion_processes(model, train, test):
    '''
    Inputs:
        model: selecting specific model
        train: df_train dataset(train_split)
        test: df_test dataset(test_split)
    Output:
        accuracy: train/test data set accuracy
        f1-score: train/test dataset f1-score
    '''
    #-----TRAIN SET-----
```

```

start = time.time()
# fitting training dataset with the algorithm
fitted_model = model.fit(train)
end = time.time()
print(f'Train_df fitting time: {end-start} seconds.')
#----- prediction on test set -----
start = time.time()
pred_train = fitted_model.transform(train)
end = time.time()
print(f'Train_df prediction time: {end-start} seconds')

train_accuracy = accuracy_evaluator.evaluate(pred_train)
train_f1_score = f1_evaluator.evaluate(pred_train)
print("Train_df accuracy rate: ", (train_accuracy))
print("Train_df f1_score:", (train_f1_score))
#-----TEST DATASET -----
start = time.time()
pred_test = fitted_model.transform(test)
end = time.time()
print("\n")
print(f'Test_df prediction time: {end-start} seconds')
# calculating accuracy and f1-score
test_accuracy = accuracy_evaluator.evaluate(pred_test)
test_f1_score = f1_evaluator.evaluate(pred_test)
print("Test_df accuracy rate: ", (test_accuracy))
print("Test_df f1_score:", (test_f1_score))

```

2.1.1 1. LOGISTIC REGRESSION CLASSIFIER NON-PARAMETRIC

```

[196]: print('LogisticRegression "Train/Test" fitting and predicting performance: \n')
model_evalutaion_processes(clf_LR, df_train, df_test)

```

LogisticRegression Train/Test fitting and predicting performance:

Train_df fitting time: 333.7577645778656 seconds.
 Train_df prediction time: 0.04687356948852539 seconds
 Train_df accuracy rate: 0.8160919540229885
 Train_df f1_score: 0.816237912789637

Test_df prediction time: 0.04390406608581543 seconds
 Test_df accuracy rate: 0.7058823529411765
 Test_df f1_score: 0.7101917690152983

2.1.2 2. LINEAR SUPPORT VECTOR CLASSIFIER MODEL:

```
[197]: print('Linear Support Vector Classifier "Train/Test" fitting and predicting_
↳performance: \n')
model_evalutaion_processes(clf_SVC, df_train, df_test)
```

Linear Support Vector Classifier "Train/Test" fitting and predicting performance:

Train_df fitting time: 785.255295753479 seconds.
Train_df prediction time: 0.03291153907775879 seconds
Train_df accuracy rate: 0.5172413793103449
Train_df f1_score: 0.5163426265590609

Test_df prediction time: 0.03137516975402832 seconds
Test_df accuracy rate: 0.5294117647058824
Test_df f1_score: 0.5359477124183007

2.1.3 3. RANDOM FOREST CLASSIFIER:

```
[198]: print('Random Forest "Train/Test" fitting and predicting performance: \n')
model_evalutaion_processes(clf_RF, df_train, df_test)
```

Random Forest "Train/Test" fitting and predicting performance:

Train_df fitting time: 847.7167658805847 seconds.
Train_df prediction time: 0.06088614463806152 seconds
Train_df accuracy rate: 0.9310344827586207
Train_df f1_score: 0.9310892172961138

Test_df prediction time: 0.047380685806274414 seconds
Test_df accuracy rate: 0.5294117647058824
Test_df f1_score: 0.5359477124183007

2.1.4 4. DECISION TREE CLASSIFIER:

```
[199]: print('Decision Tree "Train/Test" fitting and predicting performance: \n')
model_evalutaion_processes(clf_DT, df_train, df_test)
```

Decision Tree "Train/Test" fitting and predicting performance:

Train_df fitting time: 793.8072504997253 seconds.
Train_df prediction time: 0.04188823699951172 seconds
Train_df accuracy rate: 0.9310344827586207

Train_df f1_score: 0.9309798060800467

Test_df prediction time: 0.0578458309173584 seconds

Test_df accuracy rate: 0.5882352941176471

Test_df f1_score: 0.5969040247678018

PARAMETRIC CLASSIFICATION EVALUATION FUNCTION

```
[201]: def model_classification_scores(inputResult):
        """
        Prints classification scores given tp, tn, fp and fn.

        :Inputs:
            :tp: True positives
            :tn: True Negatives
            :fp: False Positives
            :fn: False Negatives

        :Print outputs:
            :accuracy: Number of correct classifications
            :precision: Number of true positives out of positive classifications
            :recall: Number of true positives out of those that should have been_
→true positives
            :f1: Harmonic mean of precision and recall, good overall stat
        """
        tp = inputResult.filter("label = 1 and prediction = 1").count()
        fp = inputResult.filter("label = 0 and prediction = 1").count()
        fn = inputResult.filter("label = 1 and prediction = 0").count()
        tn = inputResult.filter("label = 0 and prediction = 0").count()

        accuracy = (tp+tn)/(tp+tn+fp+fn)
        precision = tp / (tp + fp)
        recall = tp / (tp + fn)

        #f1 = 2 * ( precision * recall ) / (precision+recall)

        print("=====")
        print("Model evaluation metric:")
        print("\t test_Set Accuracy", accuracy)
        print("\t test_Set precision: ", precision)
        print("\t test_Set Recall: ", recall)
        # print("\t F1: ", f1)
        print("=====")

        print("Confusion matrix")
        cm = np.array([[tp, fp],
                        [fn, tn]])
```

```

columns = ['Churned', 'Remained']
rows = ['Churned', 'Remained']
# converting into a data frame
conf_df = pd.DataFrame(cm, rows, columns)
print(conf_df.head())

```

2.2 Parameterization of models:

In pursuing a better accuracy/f1 score with the train/test dataset, I think it would be beneficial to train our final model using K-Fold cross validation, which is automatically done with the CrossValidator along with a Grid Search using ParamGridBuilder.

1. LOGISTIC REGRESSION CLASSIFIER:

```

[202]: paramGrid = ParamGridBuilder().\
        addGrid(clf_LR.elasticNetParam,[0.0, 0.4, 0.8]).\
        addGrid(clf_LR.regParam,[0.00, 0.08]).\
        build()
# cross validation with parameter grid
crossval_LR = CrossValidator(estimator=clf_LR,
                             estimatorParamMaps=paramGrid,
                             evaluator=MulticlassClassificationEvaluator
                             (metricName="f1"),
                             numFolds=3)

```

```

[203]: start = time.time()
crossval_Log_Reg = crossval_LR.fit(df_train)
end = time.time()
print(f'Logistics Regression tuning with "training_dataset" is done, spent_
↳ {end-start} seconds.')

```

Logistics Regression tuning with "training_dataset" is done, spent 2914.332457780838 seconds.

```

[204]: # evaluate list metrics
crossval_Log_Reg.avgMetrics

```

```

[204]: [0.6629510584074478,
        0.7502068320205576,
        0.6629510584074478,
        0.7166851466172394,
        0.6629510584074478,
        0.7263665849441072]

```

TEST SET EVALUATION and CONFUSION MATRIX


```
[205]: # transforming test dataset
LR_pred_test = crossval_Log_Reg.transform(df_test)

# displaying a transformed test set performance
LR_pred_test.show(3)
```

```
+-----+-----+-----+-----+-----+
+
|label|          features|          rawPrediction|
probability|prediction|
+-----+-----+-----+-----+-----+
+
|    1| [-1.0762390990345...| [-0.4905910549991...| [0.37975434009131...|
1.0|
|    1| [-1.0762390990345...| [-3.3548672359297...| [0.03373614135860...|
1.0|
|    1| [0.92503195288915...| [-0.1327089789221...| [0.46687136189919...|
1.0|
+-----+-----+-----+-----+-----+
+
only showing top 3 rows
```

```
[206]: print('Test_set Accuracy with Logistic Regression: ', accuracy_evaluator.
↪evaluate(LR_pred_test))
print('Test_set F1-score with Logistic Regression: ', f1_evaluator.
↪evaluate(LR_pred_test))
```

```
Test_set Accuracy with Logistic Regression:  0.8235294117647058
Test_set F1-score with Logistic Regression:  0.826115061409179
```

```
[207]: # displaying Accuracy, Precision, Recall rates
model_classification_scores(LR_pred_test)
```

```
=====
Model evaluation metric:
      test_Set Accuracy 0.8235294117647058
      test_Set precision:  0.7142857142857143
      test_Set Recall:   0.8333333333333334
=====
```

```
Confusion matrix
      Churned  Remained
Churned      5         2
Remained      1         9
```

Confusion Maxtrix analysis:

- We have total 17 clients(test set) out of 52 total
- The classifier predicted 6 churned and 11 remained.

- In actuality we have 7 churned and 10 user remained.
- type I error 2 and type II error is 1.
- We have 2 False positive and 1 False negative, not a bad prediction?
- Since we have less cost associated with False positive, I think it is a better prediction.

Here recall(0.833) rate is higher than precision(0.714) means higher number of False Positive(2) values. Since our cost associated with False-Negative(1) is not great, I think Log Reg offers a better prediction scenario. Also our accuracy rate is 0.823 is an improvement.

FEATURE IMPORTANCE WITH LOG_REG PARAMETRIC MODEL

```
[211]: # empty dictionary
crossval_LogReg_feature_importance = {}
# Logistic regression model coefficient in array
coefficient_importance = crossval_Log_Reg.bestModel.coefficients.toArray()
# returns indices of the feature arrays
feature_coefficient_importance = np.argsort(coefficient_importance)

#-----
# looping through the coefficient arrays
for i in feature_coefficient_importance:
    crossval_LogReg_feature_importance[int(i)] = coefficient_importance[int(i)]

print("Features decreasing order of importance: " )
sorted(crossval_LogReg_feature_importance.values(), reverse=True)
```

Features decreasing order of importance:

```
[211]: [0.3779093241041484,
0.2454184685577874,
0.16752111310400472,
0.06550771637685812,
-0.01901784605560118,
-0.02601750359818768,
-0.051048411076039725,
-0.05514867311068352,
-0.061224243570619105,
-0.061644097428011746,
-0.062028010870187875,
-0.14697102232023967,
-0.24716542025503999,
-0.25592833453254066,
-0.6972170473546478]
```

```
[256]: # finding the range of coefficients of the log-reg model converting them to
↪array
coeff_ratings = crossval_Log_Reg.bestModel.coefficients.toArray()

# sorting the coefficients by order
```

```

features_indices = np.argsort(coeff_ratings)

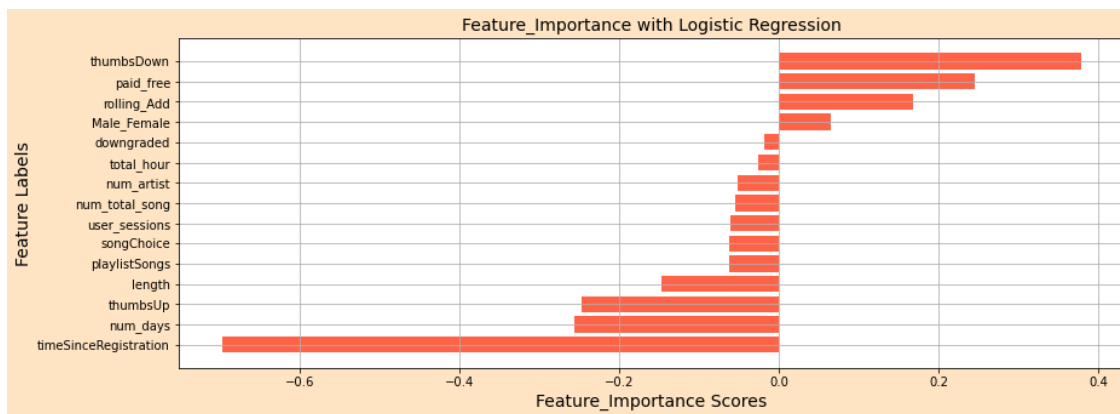
# finding features name array
final_dataSet_columns = np.array(final_dataSet_columns)

```

```

[261]: fig, ax = plt.subplots(figsize=(14, 5), edgecolor='k', facecolor='bisque')
plt.barh(range(len(features_indices)), coeff_ratings[features_indices],
         color="tomato", align="center")
plt.grid()
plt.title("Feature_Importance with Logistic Regression", fontsize=14)
plt.xlabel("Feature_Importance Scores", fontsize=14)
plt.ylabel("Feature Labels", fontsize=14)
plt.yticks(range(len(features_indices)),
           ↪final_dataSet_columns[features_indices])
plt.show()

```



2. LINEAR SUPPORT VECTOR CLASSIFICATION:

```

[215]: clf_SVC = LinearSVC()
# Grid serach of SVM
paramGrid = ParamGridBuilder().\
    addGrid(clf_SVC.maxIter, [50, 100, 700]).\
    addGrid(clf_SVC.regParam, [0.01, 0.1, 10.0, 100.0]).\
    build()

# Cross validation of SVM
Crossval_SVC = CrossValidator(estimator = clf_SVC,
                              estimatorParamMaps=paramGrid,
                              ↪
                              ↪evaluator=MulticlassClassificationEvaluator(metricName="f1"),
                              numFolds=3)

```

```
[216]: # training the model
import time
start = time.time()
model_SVC = Crossval_SVC.fit(df_train)
end = time.time()
print(f'LinearSVC() model tuning is done with "training data" total time_
↳required: {end-start} seconds.')
```

LinearSVC() model tuning is done with "training data" total time required:
25733.523796081543 seconds.

```
[217]: # list of the metrics
model_SVC.avgMetrics
```

```
[217]: [0.7276647921126765,
0.744865674480279,
0.3733802347288717,
0.3733802347288717,
0.7348286346461002,
0.744865674480279,
0.3733802347288717,
0.3733802347288717,
0.7348286346461002,
0.744865674480279,
0.3733802347288717,
0.3733802347288717]
```

```
[220]: # Hyperparameters of the best performing model
for key, value in model_SVC.getEstimatorParamMaps()[np.argmax(model_SVC.
↳avgMetrics)].items():
    print(f'{key}: {value}')
```

LinearSVC_3ce210703e76__maxIter: 50
LinearSVC_3ce210703e76__regParam: 0.1

TEST SET EVALUATION and CONFUSION MATRIX

```
[218]: SVC_pred_test = model_SVC.transform(df_test)
SVC_pred_test.show(3)
```

```
+-----+-----+-----+-----+
|label|          features|      rawPrediction|prediction|
+-----+-----+-----+-----+
|  1| [-1.0762390990345...| [-0.5711331273151...|      1.0|
|  1| [-1.0762390990345...| [-3.8924210598095...|      1.0|
|  1| [0.92503195288915...| [0.10551056733071...|      0.0|
+-----+-----+-----+-----+
only showing top 3 rows
```

```
[219]: print('Test_set Accuracy Score with Linear SVC: ', accuracy_evaluator.
        ↪evaluate(SVC_pred_test))
print('Test_set F1-Score with Linear SVC: ', f1_evaluator.
        ↪evaluate(SVC_pred_test))
```

```
Test_set Accuracy Score with Linear SVC:  0.7647058823529411
Test_set F1-Score with Linear SVC:  0.7509803921568627
```

```
[221]: model_classification_scores(SVC_pred_test)
```

```
=====
Model evaluation metric:
    test_Set Accuracy 0.7647058823529411
    test_Set precision:  0.75
    test_Set Recall:  0.5
=====
```

```
Confusion matrix
      Churned  Remained
Churned      3         1
Remained      3        10
```

CONFUSION MATRIX ANALYSIS:

- We have total 17 clients(test set) out of 52 total.
- The classifier predicted 6 churned and 11 remained.
- In actuality we have 4 churned and 13 remained.
- Our type I error(False positive) 1 and type II error(Flase negative) is 3.

Here precision(0.75) is higher than recall(0.5) rate means higher number of False Negative(3) values. In our case cost associated with FN is more than FP rate. Having a higher FN value is not a good performance for us here. Also our accuracy rate 0.76 which is not a highly improved number.

FEATURE IMPORTANCE WITH TRAINING SET:

```
[222]: # empty dictionary
SVC_feature_importance = {}
# SVC model coefficient in array
coefficient_importance = model_SVC.bestModel.coefficients.toArray()
# returns indices of the feature arrays
feature_coefficient_importance = np.argsort(coefficient_importance)

# looping through the coefficient arrays
for k in feature_coefficient_importance:
    SVC_feature_importance[int(k)] = coefficient_importance[int(k)]


print("Features in decreasing order of importance: " )
sorted(SVC_feature_importance.values(), reverse=True)
```

Features in decreasing order of importance:

```
[222]: [0.5735306469003342,  
        0.29491594953040784,  
        0.12754219796374158,  
        0.023083983094477806,  
        -0.003983469252708626,  
        -0.026042472101312036,  
        -0.03225610397048517,  
        -0.04096294714004045,  
        -0.043172327846625835,  
        -0.046569067646519934,  
        -0.09524179930665513,  
        -0.17976811498897283,  
        -0.32516465558663565,  
        -0.3735226545361331,  
        -0.8475565288222667]
```

```
[251]: # finding the range of coefficients of the log-reg model converting them to  
        ↪array  
coeff_ratings = model_SVC.bestModel.coefficients.toArray()  
  
# sorting the coefficients by order  
features_indices = np.argsort(coeff_ratings)  
  
# finding features name array  
final_dataSet_columns = np.array(final_dataSet_columns)
```

```
[262]: fig, ax = plt.subplots(figsize=(14, 5), edgecolor='k', facecolor='bisque')  
plt.barh(range(len(features_indices)), coeff_ratings[features_indices],  
         color="forestgreen", align="center")  
plt.grid()  
plt.title('Feature Importance with Linear SVC model', fontsize=14)  
plt.xlabel("Feature Importance Scores", fontsize=14)  
plt.ylabel("Feature Labels", fontsize=14)  
plt.yticks(range(len(features_indices)),  
           ↪final_dataSet_columns[features_indices])  
plt.show()
```



3. RANDOM FOREST CLASSIFIER:

```
[225]: # grid-search
paramGrid = ParamGridBuilder() \
    .addGrid(clf_RF.impurity,['entropy', 'gini']) \
    .addGrid(clf_RF.maxDepth,[5, 10]) \
    .addGrid(clf_RF.seed, [42]).build()
```

```
[226]: # cross-validation
crossval_RF = CrossValidator(estimator=clf_RF,
                             estimatorParamMaps=paramGrid,
                             evaluator=MulticlassClassificationEvaluator(),
                             numFolds=3)
```

```
[227]: # train
start = time.time()
RF_model = crossval_RF.fit(df_train)
end = time.time()
print(f'Random Forest Model tuning with "training data" is done, spent:␣
↪{end-start} seconds.')
```

Random Forest Model tuning with "training data" is done, spent:
2592.6374423503876 seconds.

```
[228]: # evaluate list of metric
RF_model.avgMetrics
```

```
[228]: [0.6078045292062572,
0.6199771669411138,
0.6940582536078905,
0.6690455513984925]
```

```
[229]: # Hyperparameters of the best performing model
for key, value in RF_model.getEstimatorParamMaps()[np.argmax(RF_model.
    ↳avgMetrics)].items():
    print(f'{key}: {value}')
```

```
RandomForestClassifier_26d4842b76e2__impurity: gini
RandomForestClassifier_26d4842b76e2__maxDepth: 5
RandomForestClassifier_26d4842b76e2__seed: 42
```

TEST SET EVALUATION and CONFUSION MATRIX

```
[230]: RF_pred_test = RF_model.transform(df_test)
RF_pred_test.show(3)
```

```
+-----+-----+-----+-----+-----+
+
|label|          features|          rawPrediction|
probability|prediction|
+-----+-----+-----+-----+-----+
+
|    1|[-1.0762390990345...| [3.96890715350467...| [0.39689071535046...|
1.0|
|    1|[-1.0762390990345...| [5.98133116883116...| [0.59813311688311...|
0.0|
|    1|[0.92503195288915...| [7.38267543859649...| [0.73826754385964...|
0.0|
+-----+-----+-----+-----+-----+
+
only showing top 3 rows
```

```
[231]: print('Test_set Accuracy score with Random Forest model: ', accuracy_evaluator.
    ↳evaluate(RF_pred_test) )
print('Test_set f1 score with Random Forest model: ', f1_evaluator.
    ↳evaluate(RF_pred_test))
```

```
Test_set Accuracy score with Random Forest model:  0.47058823529411764
Test_set f1 score with Random Forest model:  0.478345184227537
```

```
[232]: # transforming test dataset
model_classification_scores(RF_pred_test)
```

```
=====
Model evaluation metric:
    test_Set Accuracy 0.47058823529411764
    test_Set precision:  0.2857142857142857
    test_Set Recall:  0.3333333333333333
=====
Confusion matrix
```


	Churned	Remained
Churned	2	5
Remained	4	6

CONFUSION MATRIX ANALYSIS:

- We have total 17 clients(test set) out of 52 total.
- The classifier predicted 6 churned and 11 remained.
- In actuality we have 7 churned and 10 remained.
- Our type I error(False positive) 5 and type II error(Flase negative) is 4 a bigger number comparatively.

Here **precision**(0.28) is lower than **recall**(0.33) rate means higher number of **False Positive**(5) values. In our case cost associated with FN is more than FP rate. Also our accuracy rate 0.47 is way lower than other two algorithms.

FEATURE IMPORTANCE WITH TRAINING DATASET

```
[233]: RF_feature_importance = {}
feature_importance = RF_model.bestModel.featureImportances
for k in feature_importance.indices:
    RF_feature_importance[int(k)] = feature_importance[int(k)]
print('RF model feature importance decreasig order: ')
sorted(RF_feature_importance.values(), reverse=True)
```

RF model feature importance decreasig order:

```
[233]: [0.2267625083316572,
0.1268889641855206,
0.12668708803259324,
0.12387373168352141,
0.07312483723512596,
0.06080146416118844,
0.055552250843109455,
0.04535740976122087,
0.04531856279607863,
0.04392585436483098,
0.019921185689564906,
0.0169157607717078,
0.01679060053260769,
0.009654425145940258,
0.0084253564653326]
```


```
[265]: # finding the range of coefficients of the log-reg model converting them to
→array
features_ratings = RF_model.bestModel.featureImportances.toArray()

# sorting the coefficients by order
features_indices = np.argsort(features_ratings)
```

```
# finding features name array
final_dataSet_columns = np.array(final_dataSet_columns)
```

```
[268]: fig, ax = plt.subplots(figsize=(14, 5), edgecolor='k', facecolor='bisque')
plt.barh(range(len(features_indices)), features_ratings[features_indices],
         color="royalblue", align="center")

plt.grid()
plt.title('Feature Importance with Random Forest model', fontsize=14)
plt.xlabel("Feature Importance Scores", fontsize=14)
plt.ylabel("Feature Labels", fontsize=14)
plt.yticks(range(len(features_indices)),
           ↪final_dataSet_columns[features_indices])
plt.show()
```



```
[270]: # Grabbing the best estimator's feature_importances_
#importances = RF_model.bestModel.featureImportances.toArray()

# Grabbing the indices that would sort the feature importances according to
↪their importance rating
#indices = np.argsort(importances)

# Creating a features array
#final_dataSet_columns = np.array(final_dataSet_columns)

# Plotting
#plt.figure(figsize=(12,5))
#plt.title("Feature Importances")
#plt.barh(range(len(indices)), importances[indices],
#         color="peru", align="center")
#plt.yticks(range(len(indices)), final_dataSet_columns[indices])
#plt.show()
```

4. DECISION TREE CLASSIFIER

```
[236]: clf_DT = DecisionTreeClassifier() #(seed=40)
# Grid search of Decision tree
paramGrid = ParamGridBuilder() \
    .addGrid(clf_DT.impurity,['entropy', 'gini']) \
    .addGrid(clf_DT.maxDepth,[2, 4])\
    .addGrid(clf_DT.seed, [42])\
    .build()

# Cross validation of dt
crossval_DT = CrossValidator(estimator=clf_DT,
                             estimatorParamMaps=paramGrid,
                             evaluator=MulticlassClassificationEvaluator(),
                             numFolds=3)
```

```
[237]: # training the model
#import time
start = time.time()
model_DT = crossval_DT.fit(df_train)
end = time.time()
print(f'Model tuning is done and time spent: {end-start}s.')
```

Model tuning is done and time spent: 2474.098781108856s.

```
[238]: model_DT.avgMetrics
```

```
[238]: [0.7512236842865645,
0.6059160858269843,
0.7198103639280109,
0.6427050699247188]
```

```
[239]: # Hyperparameters of the best performing model
for key, value in model_DT.getEstimatorParamMaps()[np.argmax(model_DT.
    ↪avgMetrics)].items():
    print(f'{key}: {value}')
```

```
DecisionTreeClassifier_a4b724860baf__impurity: entropy
DecisionTreeClassifier_a4b724860baf__maxDepth: 2
DecisionTreeClassifier_a4b724860baf__seed: 42
```

TEST SET EVALUATION and CONFUSION MATRIX

```
[240]: DT_pred_test = model_DT.transform(df_test)
DT_pred_test.show(3)
```

```
+-----+-----+-----+-----+-----+
|label|          features|rawPrediction|          probability|prediction|
+-----+-----+-----+-----+-----+
|    1|[-1.0762390990345...|    [7.0,31.0]| [0.18421052631578...|          1.0|
```

```
|      1|[-1.0762390990345...|    [7.0,31.0]| [0.18421052631578...|      1.0|
|      1|[0.92503195288915...|   [34.0,15.0]| [0.69387755102040...|      0.0|
+-----+-----+-----+-----+
only showing top 3 rows
```

```
[241]: #DT_test_accuracy = accuracy_evaluator.evaluate(DT_pred_test)
#DT_test_f1 = f1_evaluator.evaluate(DT_pred_test)
print('Test_set Accuracy with Decision Tree model: ', accuracy_evaluator.
      ↪evaluate(DT_pred_test) )
print('Test_set F1-score with Decision Tree model: ', f1_evaluator.
      ↪evaluate(DT_pred_test))
```

```
Test_set Accuracy with Decision Tree model:  0.6470588235294118
Test_set F1-score with Decision Tree model:  0.6470588235294118
```

```
[242]: model_classification_scores(DT_pred_test)
```

```
=====
Model evaluation metric:
      test_Set Accuracy 0.6470588235294118
      test_Set precision:  0.5
      test_Set Recall:  0.5
=====
```

```
Confusion matrix
      Churned  Remained
Churned         3         3
Remained         3         8
```

CONFUSION MATRIX ANALYSIS:

- We have total 17 clients(test set) out of 52 total.
- The classifier predicted 6 churned and 11 remained.
- In actuality we have 6 churned and 11 remained.
- Our type I error(False positive) 3 and type II error(Flase negative) is 3 is equal number.

Here precision(0.50) is equal to recall(0.50) rate and evidently our False Positive(3) is equal to False Negative(3) values. Our accuracy rate is 0.64 and overall performance is slightly improved than Random Forest mdoel but not worthy of persuasion.

FEATURE IMPORTANCE WITH TRAINING SET

```
[243]: model_DT.bestModel.featureImportances
```

```
[243]: SparseVector(15, {3: 1.0})
```


```
[244]: # finding the range of coefficients of the log-reg model converting them to
      ↪array
feature_ratings = model_DT.bestModel.featureImportances.toArray()
```

```
# sorting the coefficients by order
features_indices = np.argsort(feature_ratings)

# finding features name array
final_dataSet_columns = np.array(final_dataSet_columns)
```

```
[245]: fig, ax = plt.subplots(figsize=(14, 6), edgecolor='k', facecolor='beige')
plt.barh(range(len(features_indices)), feature_ratings[features_indices],
         color="forestgreen", align="center")

plt.title('feature importance with Decision Tree model', fontsize=14)
plt.xlabel("Feature importance on a scale", fontsize=14)
plt.ylabel("Important feature labels", fontsize=14)
plt.yticks(range(len(features_indices)), □
           ↪ final_dataSet_columns[features_indices])
plt.show()
```



2.2.1 Conclusion:

It's been a journey with PySpark as I was using my desktop doing parametric computation. It's been a time ridden experience overall and I'd say using 'PySpark' intensive computation with local machine is a terrible idea. I had to navigate github, google search in extraordinary way to come this end of this project, an experience!

- I found Logistic Regression and Linear Support Vector Classification algorithm performs best among all the algorithms I've tried. But timing with LSVC is higher than Logistic regression with minimum performance gain.
- In all considerations Logistic Regression is a better choice algorithm for us here.
- The feature importance result was almost identical in both algorithms.

- On feature importance I think negatively correlated features hold importance for the ‘Not Churned’ users mostly.
- Feature engineering aspect has immense possibility of redefining the model performance.

Type I and II errors: Notes - Type I error (False Positive) refers to non-acceptance of hypothesis which ought to be accepted. By and large, Type I error crops up when the researcher notice some difference, when in fact, there is none. **So we should accept the false-positive values.**

- Type II error (False Negative) is the acceptance of hypothesis which ought to be rejected. whereas type II error arises when the researcher does not discover any difference when in truth there is one. **We should reject the false-negative values.**
- These two errors cannot be removed completely but can be reduced to a certain level.

2.3 References:

- 1: <https://thispointer.com/how-to-find-drop-duplicate-columns-in-a-dataframe-python-pandas/>
- 2.: <https://github.com/CapAllen/Sparkify/blob/master/Sparkify.ipynb>
- 3: https://github.com/ustcdj/Sparkify_Churn_Analysis/blob/master/Sparkify.ipynb
- 4: <https://www.dummies.com/education/math/statistics/how-to-set-up-a-hypothesis-test-null-versus-alternative/>

[]: